

Kernel Attacks through User- Mode Callbacks

Tarjei Mandt
SyScan'11 Taipei

Who am I

- **Security Researcher at Norman**
 - Malware Detection Team (MDT)
- **Interests**
 - Vulnerability research
 - Operating system internals
- **Past Work**
 - Kernel Pool Exploitation on Windows 7
 - Mitigating NULL Pointer Exploitation on Windows

About this Talk

- Several vulnerability classes related to windows hooks and user-mode callbacks
 - Null pointer dereferences
 - Use-after-frees
- Resulted in 44 patched privilege escalation vulnerabilities in MS11-034 and MS11-054
 - Several unannounced vulnerabilities were also addressed as part of the variant discovery process
- Requires understanding of several mechanisms specific to NT and win32k

Agenda

- Introduction
- Win32k
 - Window Manager
 - User-Mode Callbacks
- Vulnerabilities
- Exploitability
- Mitigations
- Conclusion

Introduction

Win32k and User-Mode Callbacks

Win32k

- The Windows GUI subsystem was traditionally implemented in user-mode
 - Used a client-server process model
- In NT 4.0, a large part of the server component (in CSRSS) was moved to kernel-mode
 - Introduced Win32k.sys
- Today, Win32k manages both the Window Manager (USER) and the Graphics Device Interface (GDI)

User-Mode Callbacks

- Allows win32k to make calls back into user-mode and operate on user-mode data
 - Invoke application defined hooks
 - Provide event notifications
 - Read and set properties in user-mode structures
- Implemented in the NT executive
 - nt!KeUserModeCallback
 - Works like a reverse system call

Win32k vs. User-Mode Callbacks

- Win32k uses a global locking design in creating a thread-safe environment
 - Presumably remnants of the old subsystem design
- Callbacks “interrupt” kernel execution and allow win32k structures and object properties to be modified
- Insufficient checks or validation may result in numerous vulnerabilities
 - Use-after-frees
 - NULL pointer dereferences
 - ++

Previous Work

- **Mxatone - Analyzing local privilege escalations in win32k (Uninformed vol.10)**
 - Insufficient validation of data returned from user-mode callbacks
- **Win32k Window Creation Vulnerabilities**
 - CVE-2010-0484 (MS10-032)
 - Window parent not revalidated after callbacks
 - CVE-2010-1897 (MS10-048)
 - Pseudo handle provided in callback not sufficiently validated
- **Stefan Esser - State of the Art Post Exploitation in Hardened PHP Environments (Black Hat USA 2009)**
 - Interruption vulnerabilities

Goals

- Show how user-mode callbacks without very stringent checks may introduce several subtle vulnerabilities
- Show how such vulnerabilities may be exploited using pool and kernel heap manipulation
- Propose a method to generically mitigate exploitability of NULL pointer dereference vulnerabilities

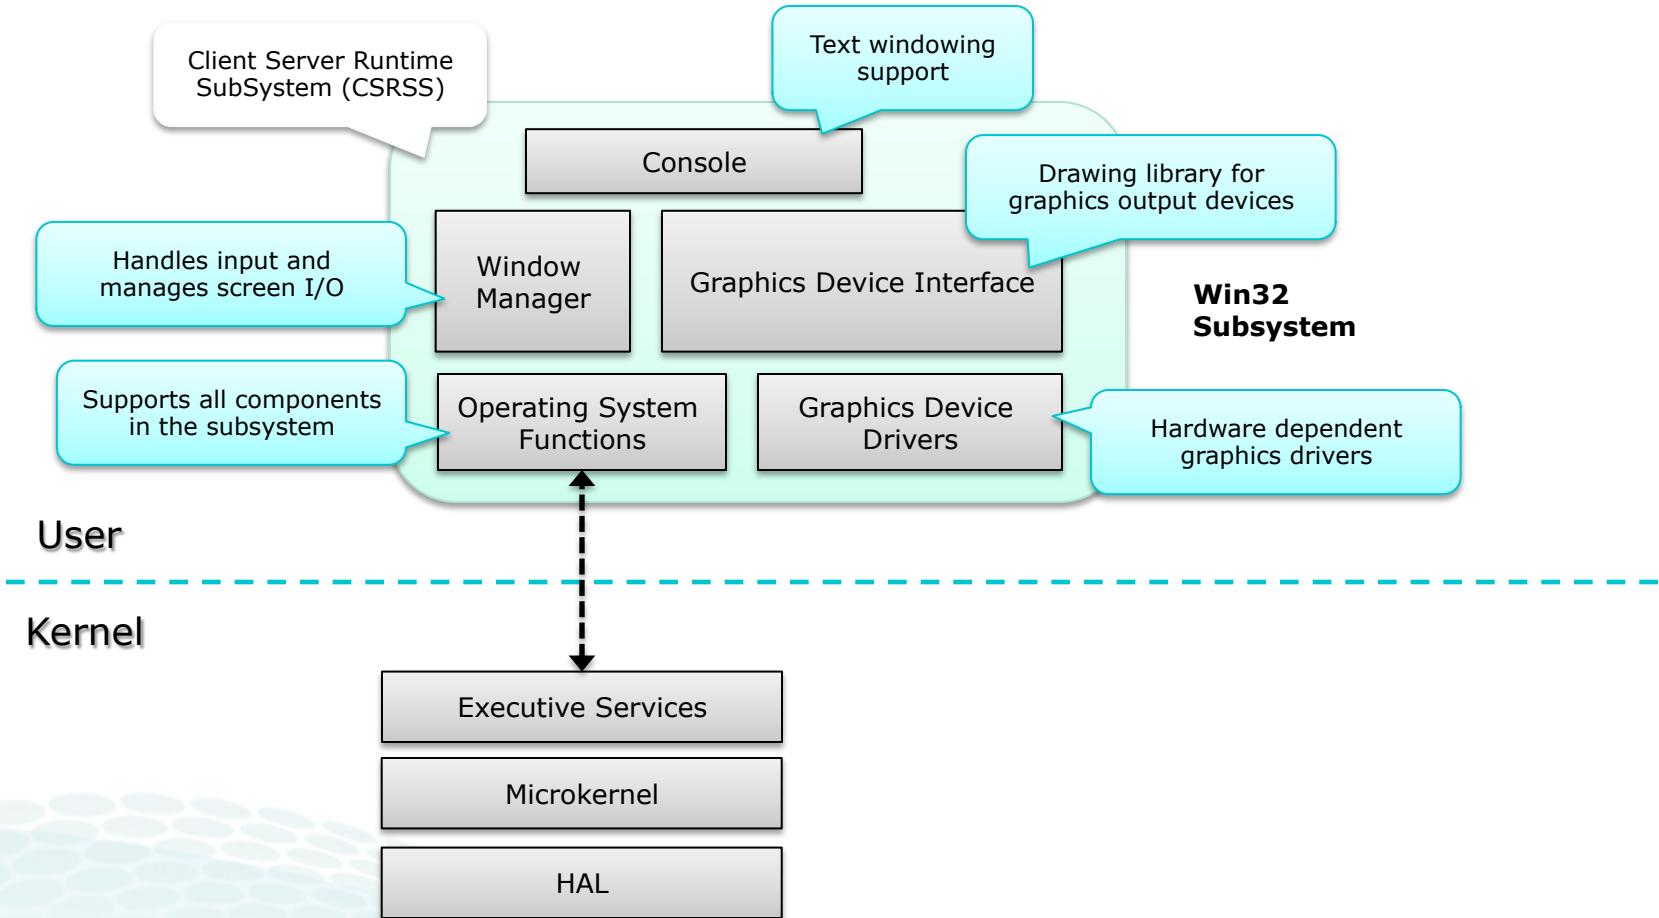
Win32k

Architecture and Design

Windows NT 3.51

- **Modified microkernel design**
 - File systems, network protocols, IPC, and drivers are implemented in kernel mode
- **Followed a more pure microkernel approach in its implementation of the GUI subsystem**
 - Window Manager and GDI implemented in the Client-Server Runtime SubSystem (CSRSS)
- **Optimized for performance**
 - Shared memory design
 - Paired threads between client and server (FastLPC)

Windows NT 3.51 Win32 Subsystem



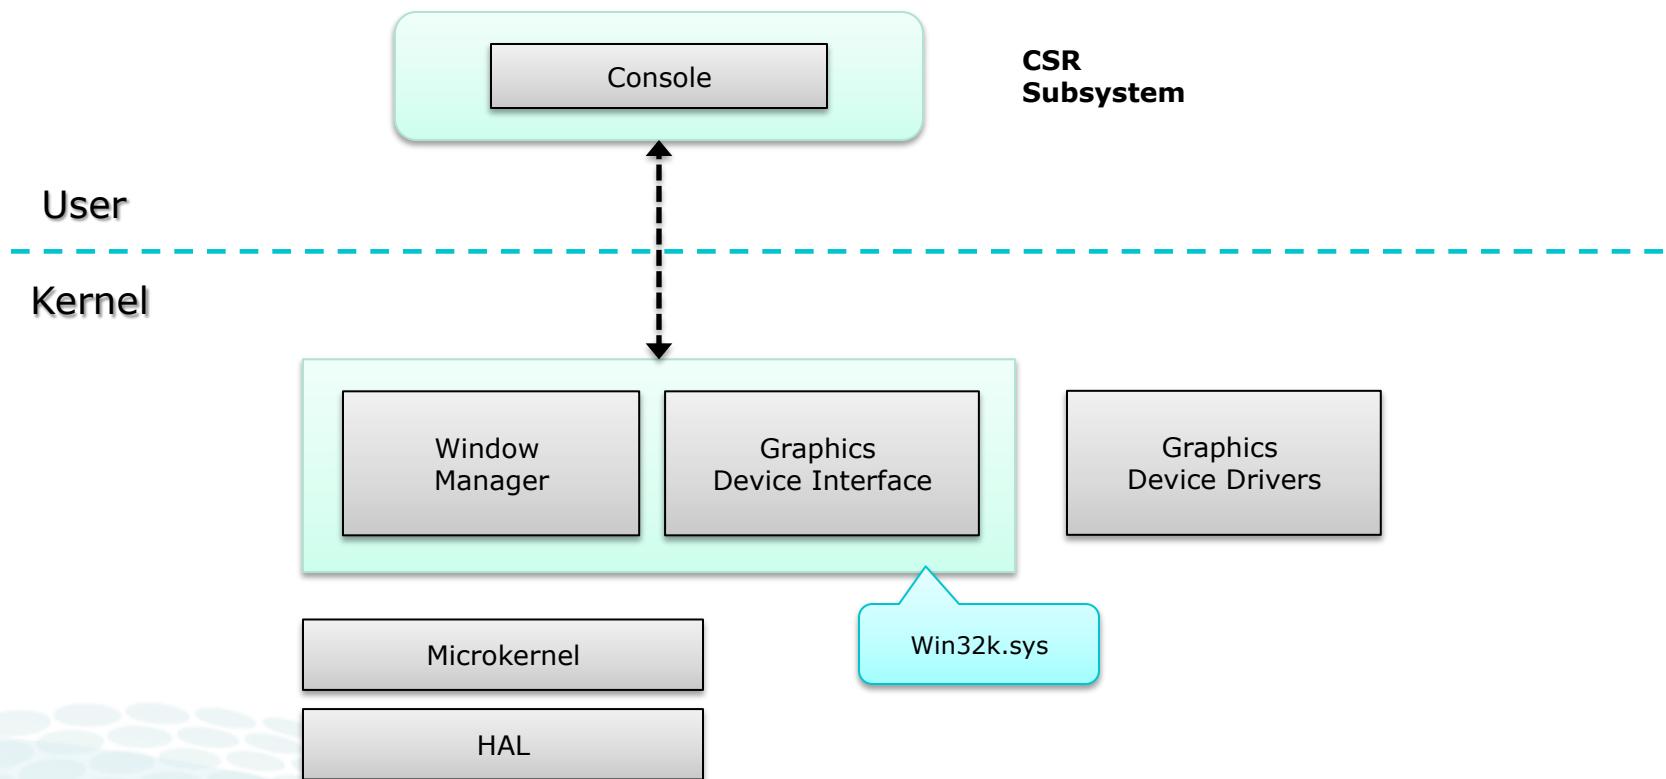
Drawbacks of the NT 3.51 Design

- **Graphics and windowing subsystem have a very high rate of interaction with hardware**
 - Video drivers, mouse, keyboard, etc.
- **Client-server interaction involves excessive thread and context switching**
 - Greatly affects graphics rendering performance
- **High memory requirements**
 - Uses 64K shared memory buffer to accumulate and pass parameters between the client and server

Windows NT 4.0

- Moved the Window Manager, GDI and graphics device drivers to kernel-mode
 - Introduced win32k.sys
- Eliminated the need for shared buffers and paired threads
 - Resulted in fewer thread and context switches
 - Reduced memory requirements
- Some old performance tricks were still maintained
 - E.g. caching of kernel structures in the user mode portion of the client's address space

Win32k.sys in Windows NT 4.0



Win32k

- Kernel component of the Win32 subsystem
- Implements the kernel side of
 - Window Manager (USER)
 - Graphics Device Interface (GDI)
- Provides thunks to DirectX interfaces
- Has its own system call table
 - More than 800 entries on Windows 7
 - win32k!W32pServiceTable

Window Manager (USER)

- **Several responsibilities**
 - Controls window displays
 - Manages screen output
 - Collects input from keyboard, mouse, etc.
 - Calls application-defined hooks
 - Passes user messages to applications
 - Manages user objects
- **The component this talk will focus on**

Graphics Device Interface (GDI)

- **Manages the graphics output and rendering**
 - Library of functions for graphics output devices
 - Includes functions for line, text, and figure drawing and for graphics manipulation
 - Manages GDI objects such as brushes, pens, DCs, paths, regions, etc.
 - Provides APIs for video/print drivers
- **Slow compared to Direct2D/DirectWrite**
 - Will probably be replaced at some point

Window Manager

User Objects and Thread Safety

User Objects

- All user handles for entities such as windows and cursors are backed by their own object
 - Allocated in win32k!HAllocateObject
- Each object type is defined by a unique structure
 - win32k!tagWND
 - win32k!tagCURSOR
- User objects are indexed into a dedicated handle table maintained by win32k
- Handle values are translated into object pointers using the handle manager validation APIs
 - win32k!HMValidateHandle(..)

User Object Header

- Every user object starts with a **HEAD** structure
- kd> dt win32k!_HEAD
 - +0x000 h : Ptr32 Void // handle value
 - +0x004 cLockObj : UInt4B // lock count
- The **lock count** tracks object use
 - An object is freed when the lock count reaches zero
- Additional fields are defined if the object is owned by a thread or process, or associated with a desktop
 - win32k!_THRDESKHEAD
 - win32k!_PROCDESKHEAD

User Handle Table

- All user objects are indexed into a per-session handle table
 - Initialized in win32k!Win32UserInitialize
- Pointer to the user handle table is stored in the win32k!tagSHAREDINFO structure
 - user32!gSharedInfo (Win 7) or win32k!gSharedInfo
- kd> dt win32k!tagSHAREDINFO
 - +0x000 psi : Ptr32 tagSERVERINFO
 - **+0x004 aheList** : Ptr32 _HANDLEENTRY
 - +0x008 HeEntrySize : Uint4B
 - +0x00c pDisplInfo : Ptr32 tagDISPLAYINFO
 - +0x010 ulSharedDelta : Uint4B

User Handle Table Entries

- Each entry in the user handle table is represented by a HANDLEENTRY structure
- kd> dt win32k!_HANDLEENTRY
 - +0x000 phead : Ptr32 _HEAD
 - +0x004 pOwner : Ptr32 Void
 - +0x008 bType : Uchar
 - +0x009 bFlags : Uchar
 - +0x00a wUniq : Uint2B
- Holds pointers to the object, its owner, type, flags, and a unique seed for the handle values
 - handle = handle_table_index | (wUniq << 0x10)
 - wUniq is incremented on object free

User Handle Table Entries

object	owner	bType	bFlags	wUniq
0	0	0	0	0
ff9d1d28	0	c	0	1
ffb7d498	ffb09678	1	40	1
ffb658f0	ffbbcc958	3	0	1
ff650618	ffb09678	1	0	1
ffb64918	ffbbcc958	3	0	1

Pointer to object
in kernel memory

Pointer to owner
(THREADINFO or
PROCESSINFO)

Object type (e.g.
window, cursor,
menu, etc.)

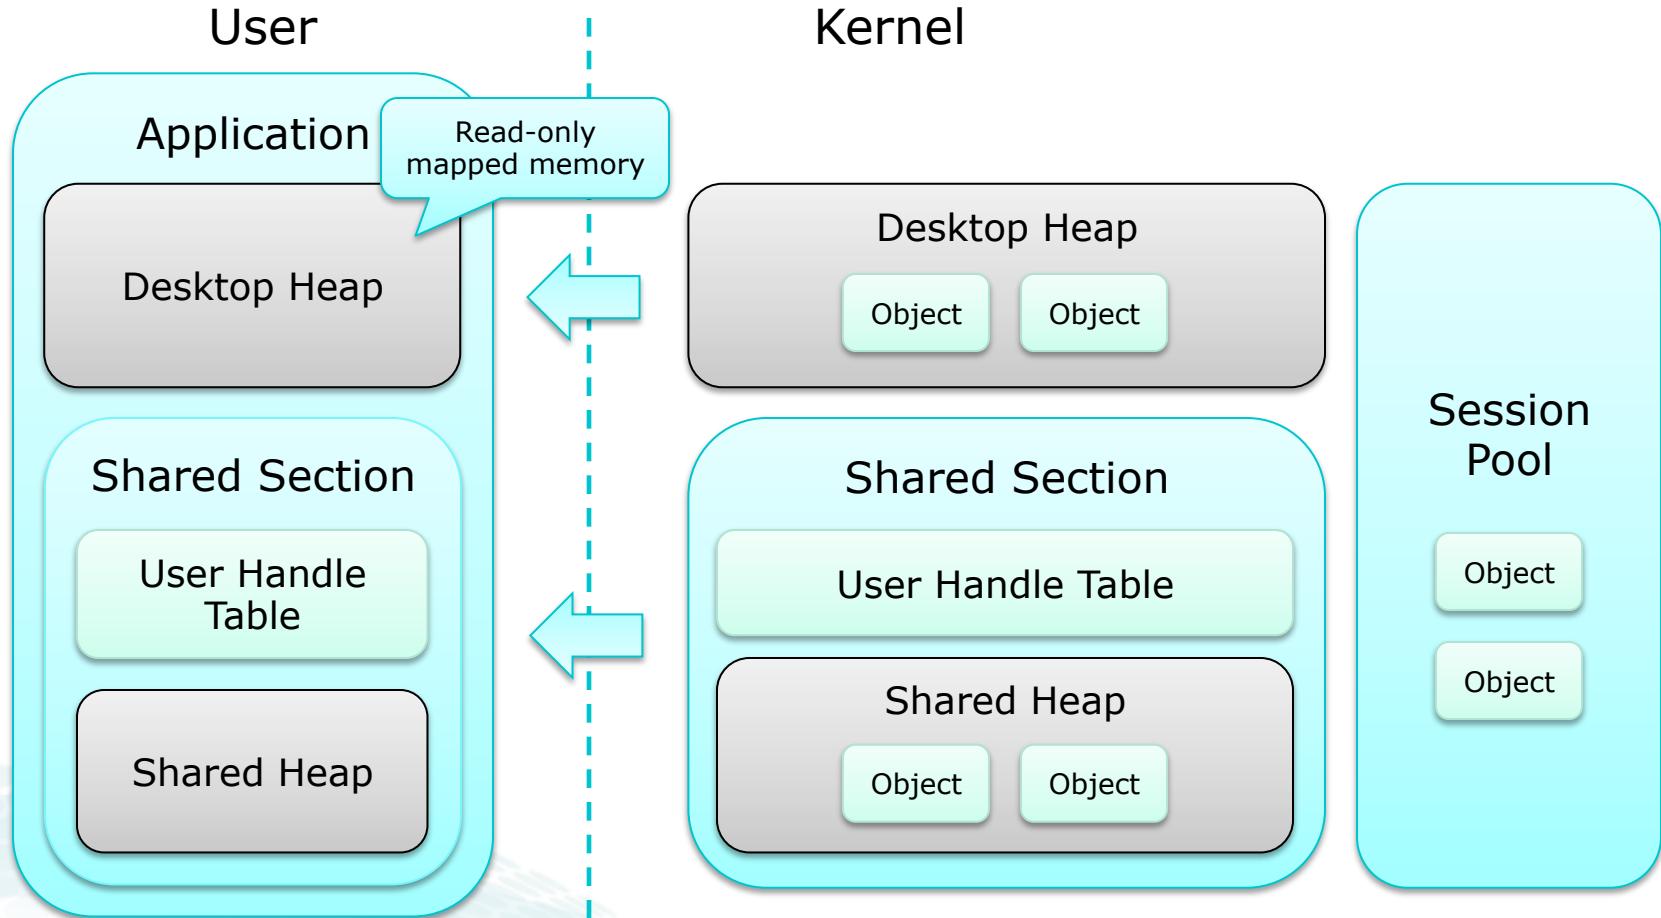
Object flags (e.g.
being destroyed)

Unique
counter

User Objects In Memory

- User objects are stored in the *session pool*, the *desktop heap* or the *shared heap*
- The desktop heap and shared heap are read-only mapped into user address space
 - Used to avoid kernel transitions
- Objects associated with a particular desktop are stored on the desktop heap
- Remaining objects are stored in the shared heap or the session pool

Handle Table & Objects In Memory



Shared Section User Mapping

- The shared section is mapped into a GUI process upon initializing the client Win32 subsystem
 - Essentially means loading user32.dll
 - Mapping itself is performed by CSRSS in calling NtUserProcessConnect (InitMapSharedSection)
- The user handle table, at the base of the shared section, can be obtained in at least two ways
 - From user32!gSharedInfo (exported on Windows 7)
 - From the connection information buffer returned by **CsrClientConnectToServer** upon specifying **USERSRV_SEVERDLL_INDEX (3)**

Handle Table From User-Mode

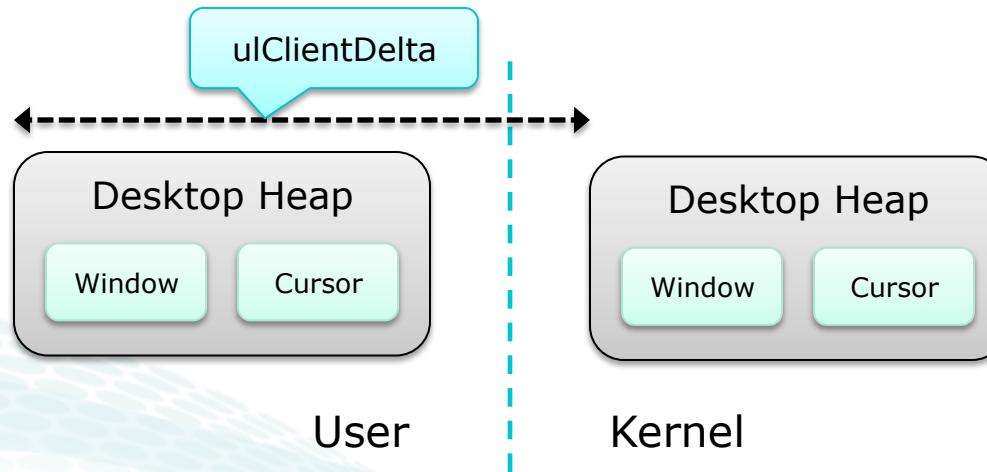
Index	Handle	Object	Owner	Type
[0000]	10000	0	0	0 <Free>
[0001]	10001	bc5d1b48	0	c <Monitor>
[0002]	10002	e1a12698	e1a13008	1 <Window>
[0003]	10003	e15a91f8	e15ad650	3 <Icon/Cursor>
[0004]	10004	bc6006e8	e1a13008	1 <Window>
[0005]	10005	e163c670	e15ad650	3 <Icon/Cursor>
[0006]	10006	bc600818	e1a13008	1 <Window>
[0007]	10007	e15aee80	e15ad650	3 <Icon/Cursor>
[0008]	10008	bc600940	e1a13008	1 <Window>
[0009]	10009	e15aee20	e15ad650	3 <Icon/Cursor>
[000a]	1000a	bc600a88	e1a13008	1 <Window>
[000b]	1000b	e15adb80	e15ad650	3 <Icon/Cursor>
[000c]	1000c	bc6206e8	e1a13008	1 <Window>
[000d]	1000d	e17c2658	e15ad650	3 <Icon/Cursor>
[000e]	1000e	bc620818	e1a13008	1 <Window>
[000f]	1000f	e17c1610	e15ad650	3 <Icon/Cursor>
[0010]	10010	bc620940	e1a13008	1 <Window>
[0011]	10011	e17b22a8	e15ad650	3 <Icon/Cursor>
[0012]	10012	bc620a88	e1a13008	1 <Window>
[0013]	10013	e17d7e20	e15ad650	3 <Icon/Cursor>
[0014]	10014	bc6306e8	e1a13008	1 <Window>
[0015]	10015	e17d7dc0	e15ad650	3 <Icon/Cursor>

Desktop Heap User Mapping

- For each GUI thread, win32k maps the associated desktop heap into the user-mode process
 - Performed by win32k!MapDesktop
- Information on the desktop heap is stored in the desktop information structure
 - Holds the kernel address of the desktop heap
 - Accessible from user-mode
 - NtCurrentTeb()->Win32ClientInfo.pDeskInfo
- **kd> dt win32k!tagDESKTOPINFO**
 - +0x000 pvDesktopBase : Ptr32 Void
 - +0x004 pvDestkopLimit : Ptr32 Void

Kernel-Mode -> User-Mode Address

- User-space address of desktop heap objects are computed using `ulClientDelta`
 - `NtCurrentTeb()->Win32ClientInfo.ulClientDelta`
- User-space address of shared heap objects are computed using `ulSharedDelta`
 - Defined in `win32k!tagSHAREDINFO`



User Object From User-Mode

```
C:\> C:\WINDOWS\system32\cmd.exe
C:\Do...ys\vmware\Desktop>sharedinfo.exe 20096 a4
[*] Dumping object data for handle: 20096
0000h: 00020096 00000001 e183f720 818d6238 .....8b..
0010h: bbe8c818 00000000 80000300 c0000800 .....T.
0020h: 54009945 00000000 00000000 00000000 E.....h.
0030h: 00000000 bbe8c768 00000000 00000000 .....|....&.
0040h: 00000000 00000004 0000017c 00000026 .....|....&.
0050h: 00000000 00000004 0000017c 00000026 .....|....&.
0060h: 723e208b bbe88328 00000000 bbe8c700 ..>w<...
0070h: 00000000 00000000 00000000 00000000 .....|.....
0080h: 00000000 00000000 00000000 00000004 .....|.....
0090h: 00000000 00000000 0009e508 .....|.....
C:\>
```

User Object Types

- On Windows 7, there are 21 different user object types (22 including the ‘free’ type)
 - Includes ‘touch’ and ‘gesture’ objects
- Information on each type is stored in the handle type information table
 - win32k!ghti (undocumented structure)
 - Defines the destroy routines for each type
 - Defines target memory location (desktop/shared heap, session pool)

User Object Types #1

ID	TYPE	OWNER	MEMORY
0	Free		
1	Window	Thread	Desktop Heap / Session Pool *
2	Menu	Process	Desktop Heap
3	Cursor	Process	Session Pool
4	SetWindowPos	Thread	Session Pool
5	Hook	Thread	Desktop Heap
6	Clipboard Data		Session Pool
7	CallProcData	Process	Desktop Heap
8	Accelerator	Process	Session Pool
9	DDE Access	Thread	Session Pool
10	DDE Conversation	Thread	Session Pool

* Stored on the desktop heap if the window is associated with a desktop

User Object Types #2

ID	TYPE	OWNER	MEMORY
11	DDE Transaction	Thread	Session Pool
12	Monitor		Shared Heap
13	Keyboard Layout		Session Pool
14	Keyboard File		Session Pool
15	Event Hook	Thread	Session Pool
16	Timer		Session Pool
17	Input Context	Thread	Desktop Heap
18	Hid Data	Thread	Session Pool
19	Device Info		Session Pool
20 (Win 7)	Touch	Thread	Session Pool
21 (Win 7)	Gesture	Thread	Session Pool

User Critical Section

- Unlike NT, the Window Manager does not exclusively lock each user object
 - Implements a global lock per session
- Each kernel routine that operates on win32k structures or objects must first acquire a lock on win32k!gpresUser
 - Exclusive lock used if write operations are involved
 - Otherwise, shared lock is used
- Clearly not designed to be multithreaded
 - E.g. two separate applications in the same session cannot process their message queues simultaneously

Shared and Exclusive Locks

The image shows two side-by-side debugger windows. The left window displays assembly code for a shared lock operation, and the right window displays assembly code for an exclusive lock operation. A callout bubble points from the shared lock code to the shared lock text, and another callout bubble points from the exclusive lock code to the exclusive lock text.

Left Window (Shared Lock):

```
; Attributes: bp-based frame
; __stdcall NtUserCheckDesktopByThreadId(x)
_NtUserCheckDesktopByThreadId@4 proc near

var_4= dword ptr -4
arg_0= dword ptr 8

mov    edi, edi
push   ebp
mov    ebp, esp
push   ecx
push   esi
call   _EnterSharedCrit@0 ; EnterSharedCrit()
mov    esi, eax
call   ds:_imp__PsGetCurrentProcess@0 ; PsGetCurrentProc
push   eax
call   _IsProcessDwm@4 ; IsProcessDwm(x)
test   eax, eax
jnz    short loc_BF844222
```

Acquire shared lock

Right Window (Exclusive Lock):

```
; Attributes: bp-based frame
; int __stdcall NtUserSwitchDesktop(HANDLE Handle, int)
_NtUserSwitchDesktop@8 proc near

var_10= byte ptr -10h
var_4= dword ptr -4
Handle= dword ptr 8
arg_4= dword ptr 0Ch

mov    edi, edi
push   ebp
mov    ebp, esp
sub   esp, 10h
push   esi
call   _UserEnterUserCritSec@0 ; UserEnterUserCritSec()
mov    eax, _gptiCurrent
xor    esi, esi
test   dword ptr [eax+00000004]
jz    short loc_BF8198
```

Acquire exclusive lock

User-Mode Callbacks

Kernel to User Interaction

User-Mode Callbacks

- In interacting with user-mode data, win32k is required to make calls back into user-mode
 - Lead to the concept of user-mode callbacks
- Implemented in `nt!KeUserModeCallback`
 - Works like a reverse system call
 - Previously researched by Ivanlef0u and mxatone, among others
- Used extensively in user object handling
 - Some user objects store data in user-mode

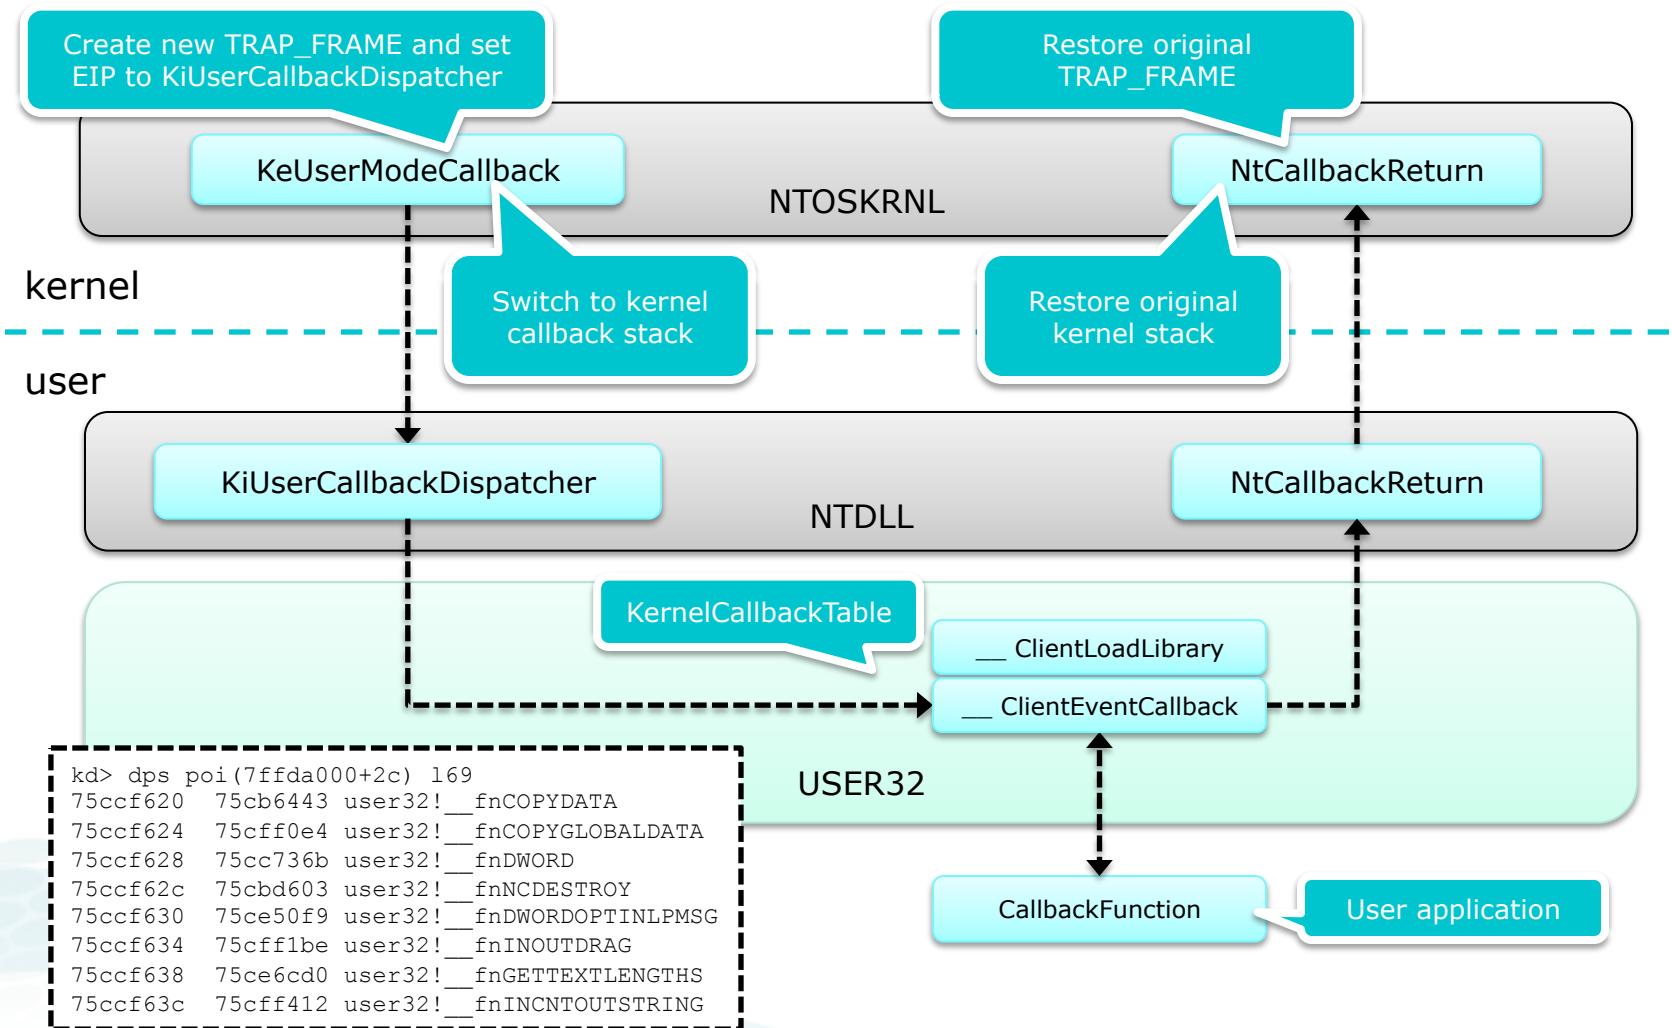
KeUserModeCallback

- **NTSTATUS KeUserModeCallback (**
 IN ULONG ApiNumber,
 IN PVOID InputBuffer,
 IN ULONG InputLength,
 OUT PVOID * OutputBuffer,
 IN PULONG OutputLength);
- ***ApiNumber* is an index into the user-mode callback function table**
 - Pointer copied to the Process Environment Block (PEB) during the initialization of USER32.dll in a process
- **kd> dt nt!_PEB KernelCallbackTable**
 - +0x02c KernelCallbackTable : Ptr32 Void

KeUserModeCallback Internals

- In a system call, a trap frame is stored on the kernel thread stack by `KiSystemService` or `KiFastCallEntry`
 - Used to save thread context and restore registers upon returning to user-mode
- `KeUserModeCallback` creates a new trap frame (`KTRAP_FRAME`) before invoking `KiServiceExit`
 - Sets EIP to `ntdll!KiUserCallbackDispatcher`
 - Replaces TrapFrame pointer of the current thread
- Input buffer is copied to the user-mode stack

KeUserModeCallback



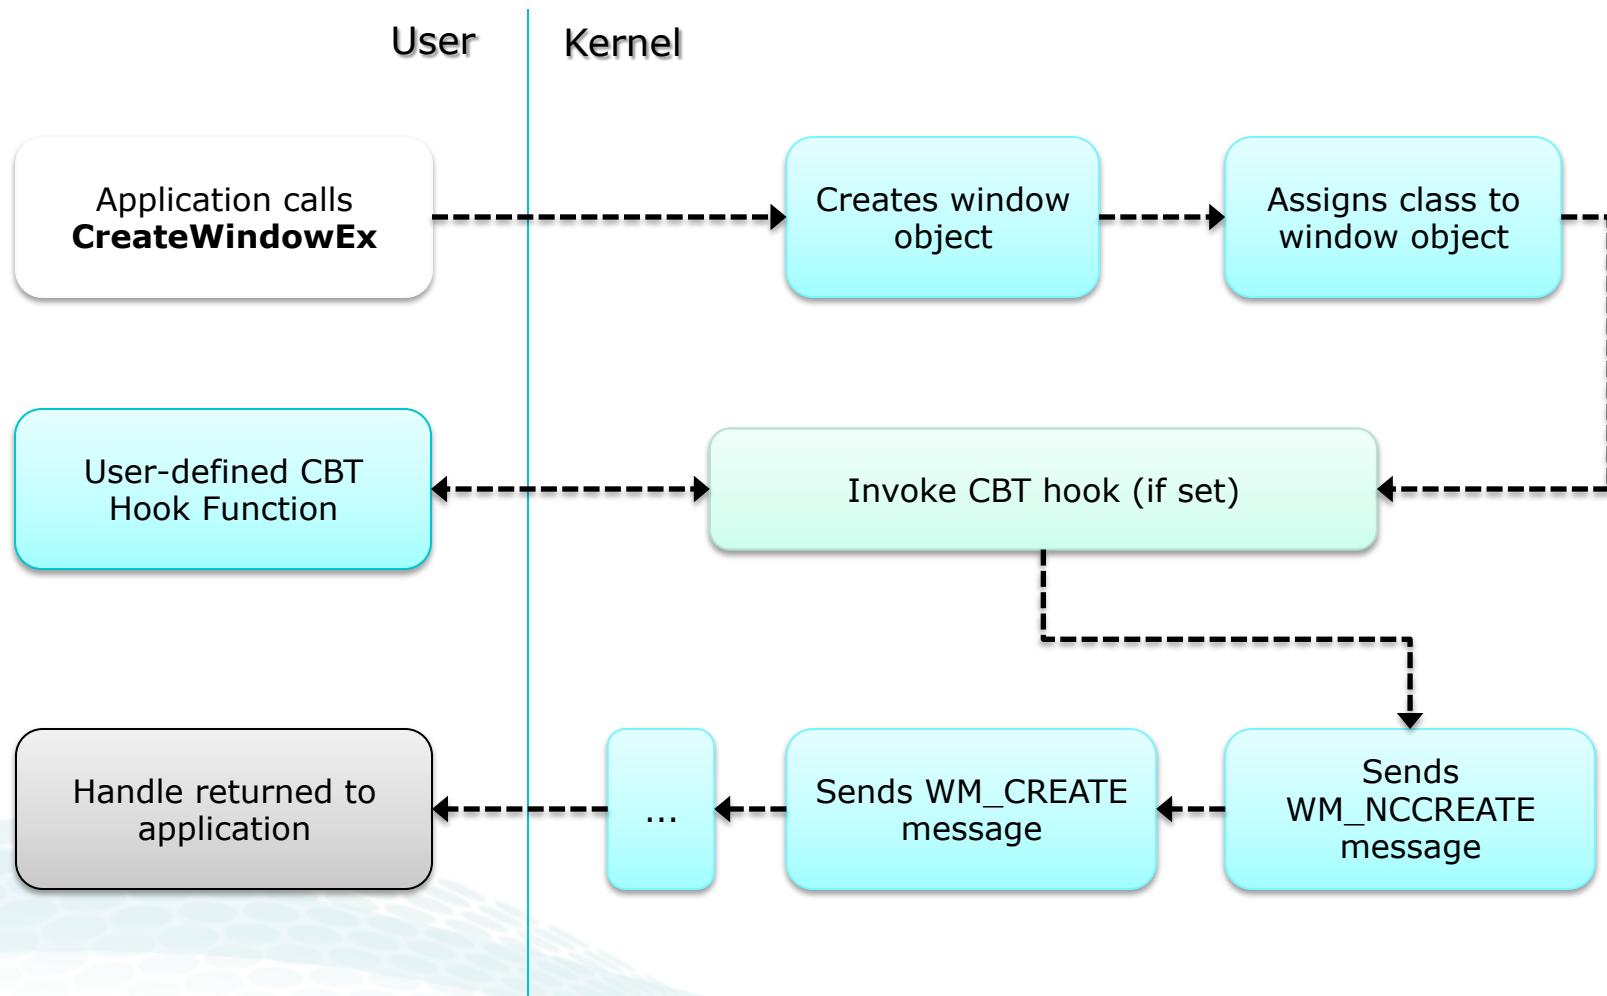
NtCallbackReturn

- **NTSTATUS NtCallbackReturn (**
 IN PVOID Result OPTIONAL,
 IN ULONG ResultLength,
 IN NTSTATUS Status);
- Used to resume execution in the kernel after a user-mode callback
- Copies the result of the callback back to the original kernel stack
- Restores original trap frame and kernel stack by using the information held in the KSTACK_AREA
- Deletes the kernel callback stack upon completion

Applications of User-Mode Callbacks

- **User-mode callbacks allow win32k to perform a variety of tasks**
 - Invoke application-specific windows hooks
 - Provide event notification
 - Copy data to and from user-mode (e.g. for DDE)
- **Hooks allow users to execute code in response to certain actions performed by win32k**
 - Calling a window procedure
 - Creating or destroying a window
 - Processing keyboard or mouse input

CreateWindow CBT Hook Example



Kernel Attacks through User-Mode Callbacks

Vulnerabilities in Win32k

User Critical Section vs. Callbacks

- Whenever a callback is executed, the kernel leaves the win32k user critical section
 - Allows win32k to perform other tasks while user-mode code is being executed
- Upon returning from a callback, win32k must ensure that referenced objects are still in the expected state
 - E.g. a callback could call SetParent() to update the parent of a window
- Insufficient checks may lead to vulnerabilities

Function Name Decoration

- Win32k.sys uses function name decoration to keep track of functions that leave the critical section
 - Prefixed “xxx” and “zzz”
- Functions prefixed “xxx” may leave the critical section and invoke a user-mode callback
 - May sometimes require a specific argument or set of arguments to trigger the actual callback
- Functions prefixed “zzz” typically invoke a deferred event callback
 - However, if win32k!gdwDeferWinEvent is null, an immediate callback is performed

Function Name Decoration Issues

- Functions that leave the critical section and invoke user-mode callbacks are not always prefixed
 - Could lead to invalid assumptions by the programmer
 - Easy to spot using IDAPython and cross referencing
- Lack of consistency in behavior of “zzz” functions
 - Some “zzz” functions seem to increment gdwDeferWinEvent while others do not

Windows 7 RTM	Windows 7 (MS11-034)
MNRecalcTabStrings	xxxMNRecalcTabStrings
FreeDDEHandle	xxxFreeDDEHandle
ClientFreeDDEHandle	xxxClientFreeDDEHandle

Locating Undecorated Functions

IDA - I:\research\kernel\win7_rtm\win32k.idb (win32k.sys)

File Edit Jump Search View Debugger Options Windows Help

IDA View-A Hex View-A Structures Enums Imports Exports

MNRecalcTabStrings(x,x,x,x,x,x) arg_10 = dword ptr 18h
MNRecalcTabStrings(x,x,x,x,x,x) arg_14 = dword ptr 1Ch
MNRecalcTabStrings(x,x,x,x,x,x)
MNRecalcTabStrings(x,x,x,x,x,x)
MNRecalcTabStrings(x,x,x,x,x,x)+2 mov edi, edi
MNRecalcTabStrings(x,x,x,x,x,x)+3 push ebp
MNRecalcTabStrings(x,x,x,x,x,x)+5
MNRecalcTabStrings(x,x,x,x,x,x)+8
MNRecalcTabStrings(x,x,x,x,x,x)+B
MNRecalcTabStrings(x,x,x,x,x,x)+E
MNRecalcTabStrings(x,x,x,x,x,x)+12
MNRecalcTabStrings(x,x,x,x,x,x)+15
MNRecalcTabStrings(x,x,x,x,x,x)+18
FUNCTION: _CreateTerminalInput
FUNCTION: _PostMessageExtended@20
FUNCTION: _QueueNotifyMessage@20 ->
FUNCTION: _NotifyOverlayWindow@20 ->
FUNCTION: _PostMessage@20 ->
FUNCTION: _QueueNotify@20 ->
FUNCTION: _DrawIconCallBack@16 ->
FUNCTION: _Notify@20 ->
FUNCTION: _DrawSwitchWndHilite@20 ->
FUNCTION: _DrawCtlThumb@4 ->
FUNCTION: _RawInputThread@4 ->
FUNCTION: _Win32kPnpNotify@4 ->
FUNCTION: _VideoPortCalloutThread@4 ->
FUNCTION: _UserInitialize@0 ->
FUNCTION: _MNRecalcTabStrings@24 ->
FUNCTION: _DT_GetExtents@20 ->
FUNCTION: _Reply@4 ->
FUNCTION: _ZapActive@4 ->
FUNCTION: _CancelDm@4 ->
FUNCTION: _Destroy@4 ->
FUNCTION: _SendDm@4 ->
FUNCTION: _UserThrd@4 ->
FUNCTION: _Client@4 ->
FUNCTION: _FreeDDE@4 ->
FUNCTION: _DT_Draw@4 ->
All done...
Python
AU: idle Down

Undecorated functions that potentially may invoke callbacks

FUNCTION: _CreateTerminalInput
FUNCTION: _PostMessageExtended@20 ->
FUNCTION: _QueueNotifyMessage@20 ->
FUNCTION: _NotifyOverlayWindow@20 ->
FUNCTION: _PostMessage@20 ->
FUNCTION: _QueueNotify@20 ->
FUNCTION: _DrawIconCallBack@16 ->
FUNCTION: _Notify@20 ->
FUNCTION: _DrawSwitchWndHilite@20 ->
FUNCTION: _DrawCtlThumb@4 ->
FUNCTION: _RawInputThread@4 ->
FUNCTION: _Win32kPnpNotify@4 ->
FUNCTION: _VideoPortCalloutThread@4 ->
FUNCTION: _UserInitialize@0 ->
FUNCTION: _MNRecalcTabStrings@24 ->
FUNCTION: _DT_GetExtents@20 ->
FUNCTION: _Reply@4 ->
FUNCTION: _ZapActive@4 ->
FUNCTION: _CancelDm@4 ->
FUNCTION: _Destroy@4 ->
FUNCTION: _SendDm@4 ->
FUNCTION: _UserThrd@4 ->
FUNCTION: _Client@4 ->
FUNCTION: _FreeDDE@4 ->
FUNCTION: _DT_DrawStr@36 ->
All done...

Search for functions that may call KeUserModeCallback or leave the user critical section

Object Locking

- Objects expected to be valid after the kernel leaves the user critical section must be *locked*
 - The *cLockObj* field of the user object header stores the object reference count
- Two forms of locking
 - Thread locking
 - Assignment locking

Thread Locking

- Used to lock objects or buffers within the context of a thread
 - ThreadLock* (inlined mostly) and ThreadUnlock*
- Each thread locked entry is stored as a TL structure in a thread-specific thread lock list
 - kd> dt win32k!_TL
 - +0x000 next : Ptr32 _TL
 - +0x004 pobj : Ptr32 Void
 - +0x008 pfnFree : Ptr32 Void
- Upon thread termination, the thread lock list is processed to release any outstanding entries
 - xxxDestroyThreadInfo -> DestroyThreadsObjects

Thread Locking By Example

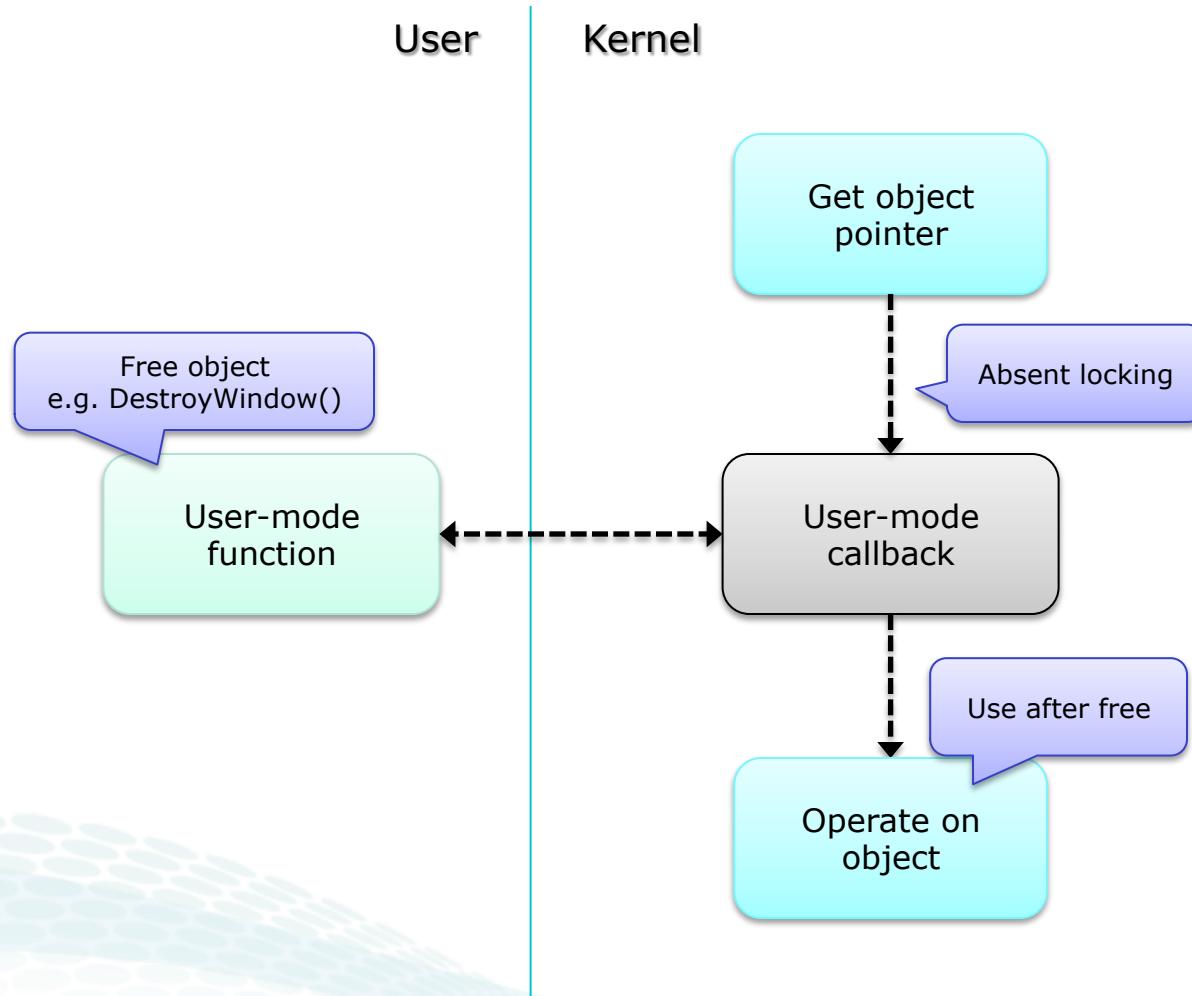
Assignment Locking

- The handle manager provides functions for thread independent locking of objects
 - HMAssignmentLock(Address, Object)
 - HMAssignmentUnlock(Address)
- Assignment locking an object to an address with an initialized pointer, releases the existing reference
- Does not provide the safety net thread locking does
 - E.g. if a thread termination occurs in a callback, the thread cleanup code must release these references

Object Locking Vulnerabilities

- Any object expected to be valid after a user-mode callback should be locked
- Similarly, any object that no longer is used by a particular component should be released
- Mismanagement in the locking and release of objects could result in the following
 - No retention: An object could be freed too early
 - No release: An object could never be freed, or the reference count (e.g. 32-bit on x86) could wrap

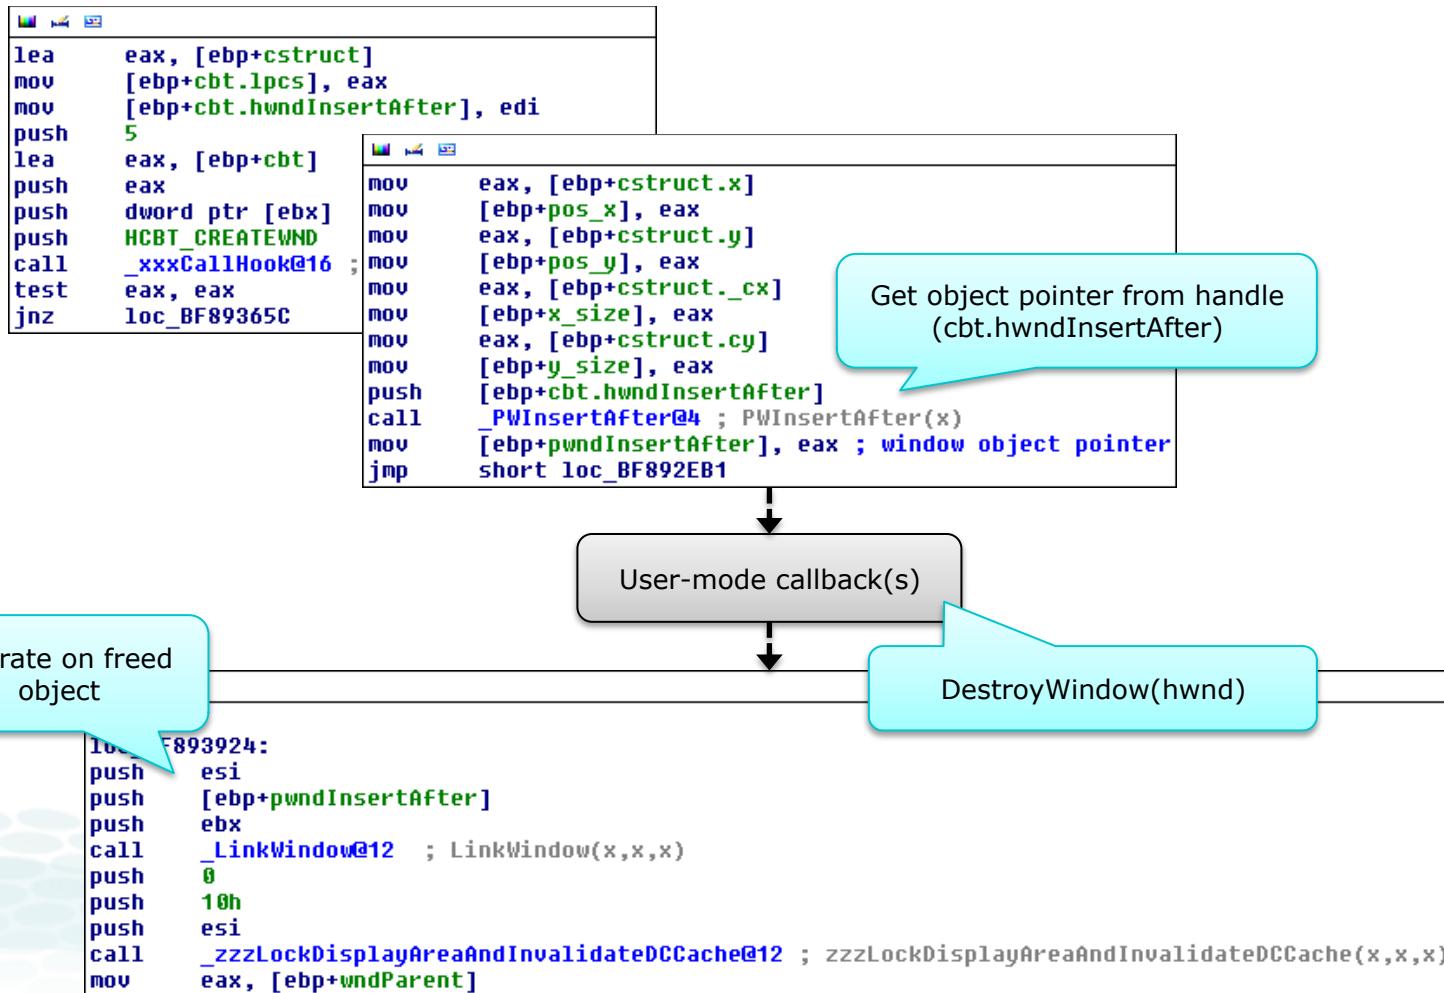
Object Use-After-Free



Window Object Use-After-Free

- In creating a window, an application can adjust its orientation and z-order using a CBT hook
 - Z-order is defined by providing the handle to the window after which the new window is inserted
- **win32k!xxxCreateWindowEx failed to properly lock the provided z-order window**
 - Only stored a pointer to the object in a local variable
- An attacker could destroy the window in a subsequent user-mode callback and trigger a use-after-free

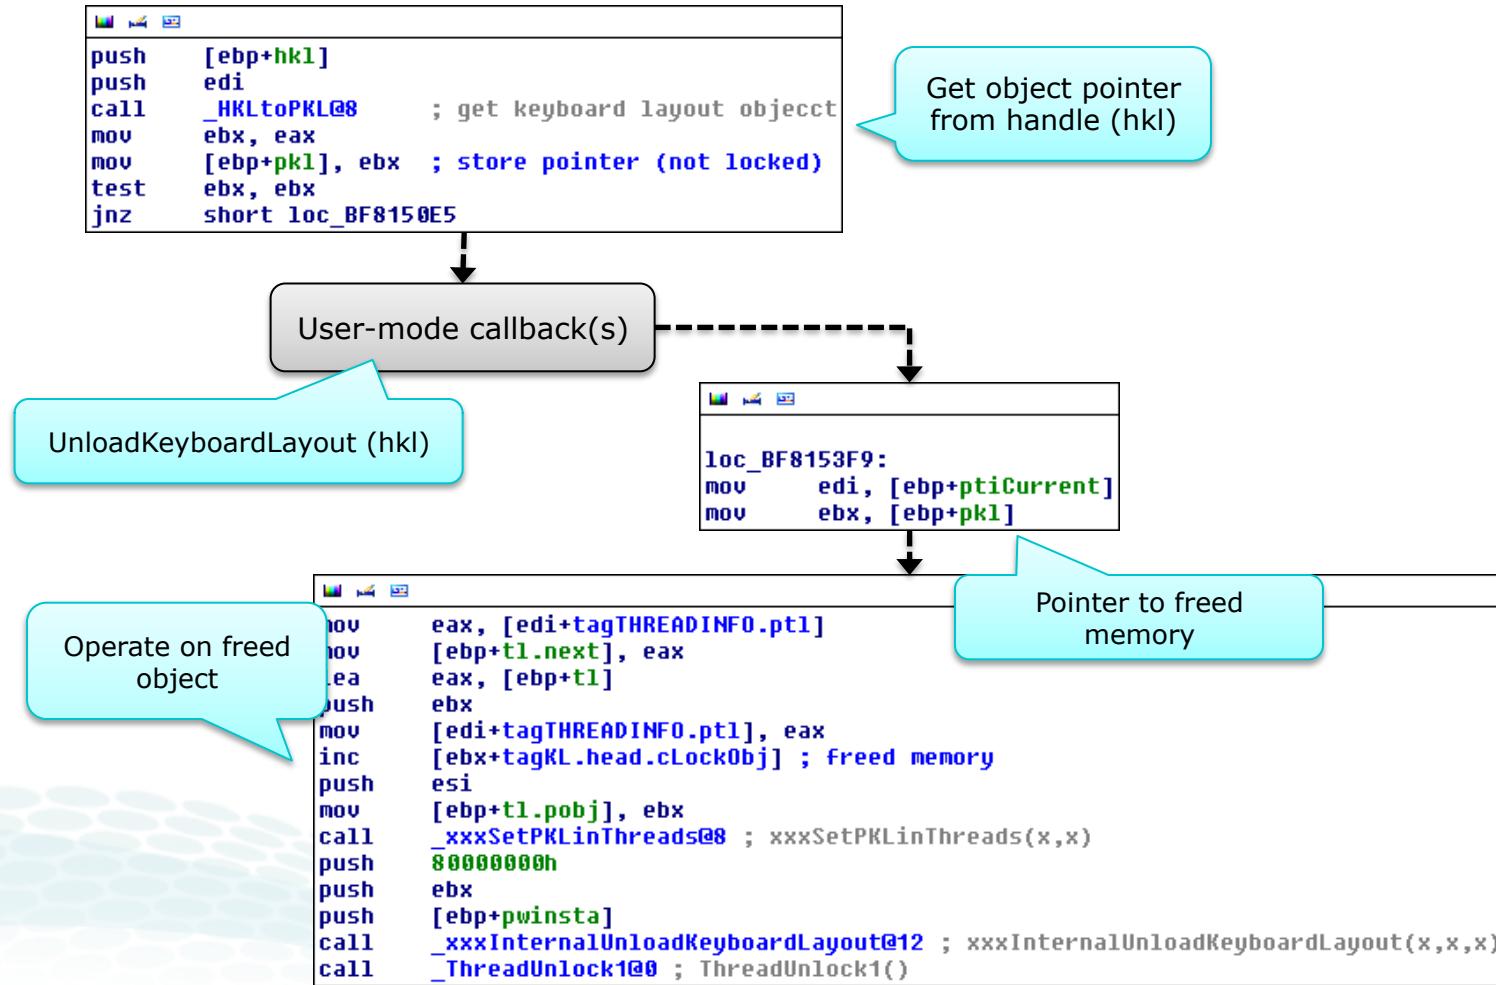
Window Object Use-After-Free



Keyboard Layout Object Use-After-Free

- In loading a keyboard layout, win32k! `xxxLoadKeyboardLayoutEx` did not lock the keyboard layout object
 - Pointer stored in local variable
- An attacker could unload the keyboard layout in a user-mode callback and thus free the object
- Subsequently, upon using the object pointer the kernel would operate on freed memory

Keyboard Layout Object Use-After-Free



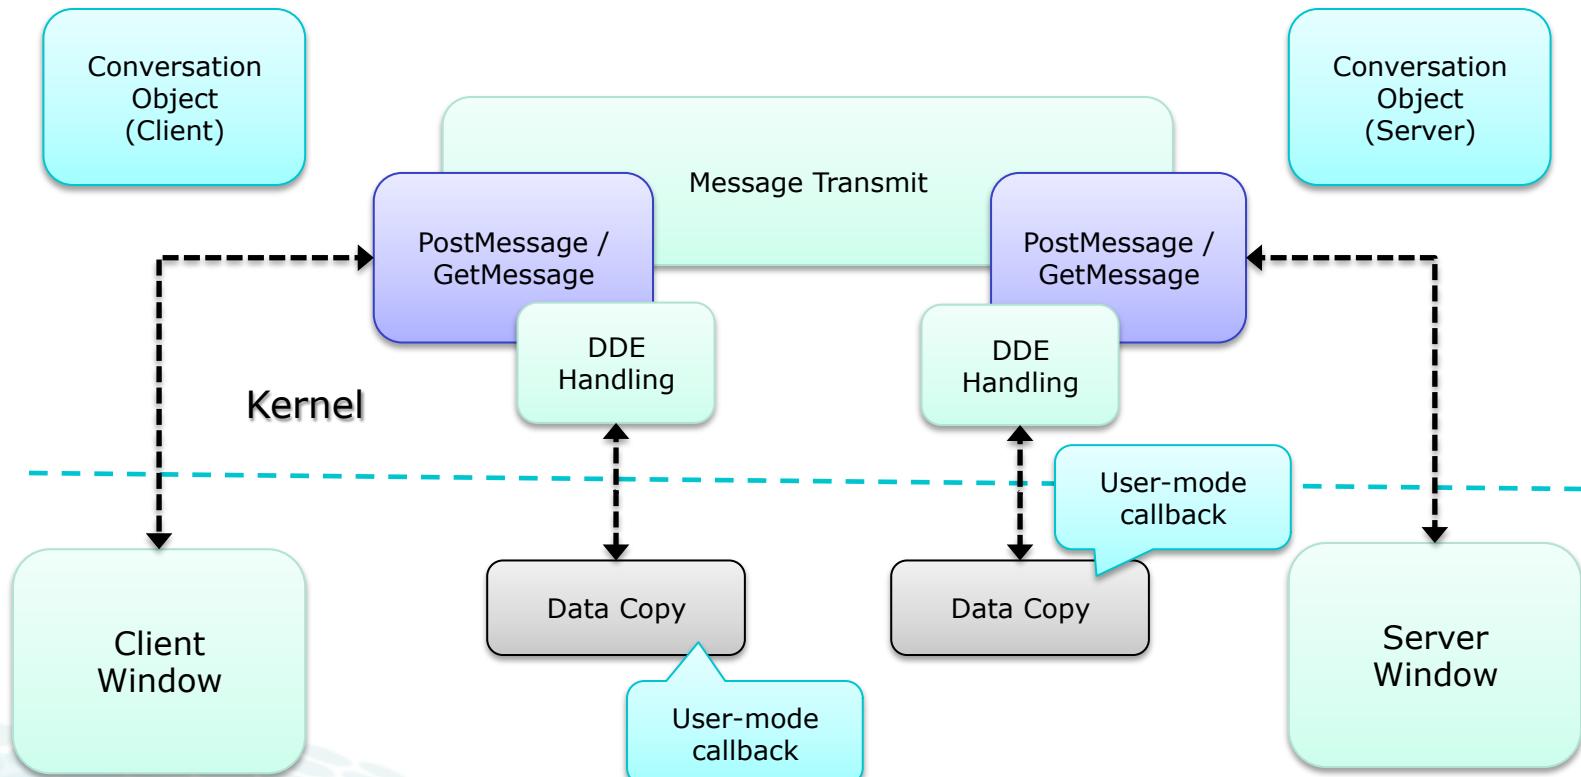
Object State Validation

- **Objects assumed to be in a certain state should always have their state validated**
 - Usually involves checking for initialized pointers or flags
- **User-mode callbacks could alter the state and update properties of objects**
 - A drop down menu is no longer active
 - The parent of a window has changed
 - The partner in a DDE conversation terminated

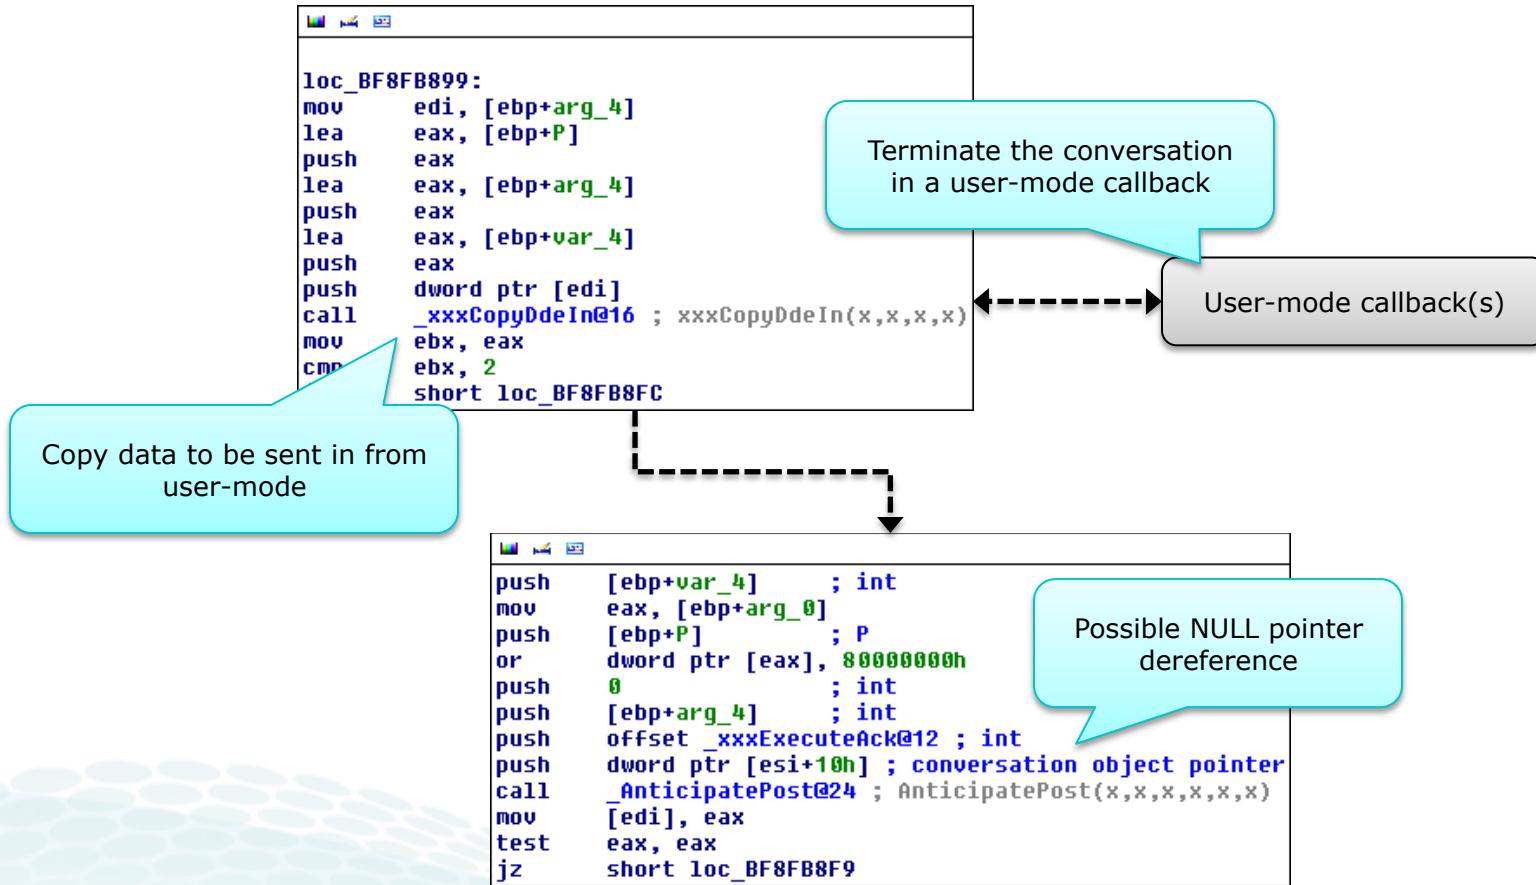
DDE Conversation State Vulnerabilities

- **Dynamic Data Exchange (DDE)**
 - Legacy protocol using messages and shared memory to exchange data between applications
- **Several functions did not sufficiently validate DDE conversation objects after user-mode callbacks**
 - Used to copy data in and out from user-mode
- **An attacker could terminate a conversation in a user-mode callback and thus unlock the partner conversation object**
 - Could result in a NULL pointer dereference as the function did not revalidate the conversation object pointer

DDE Conversation Message Handling



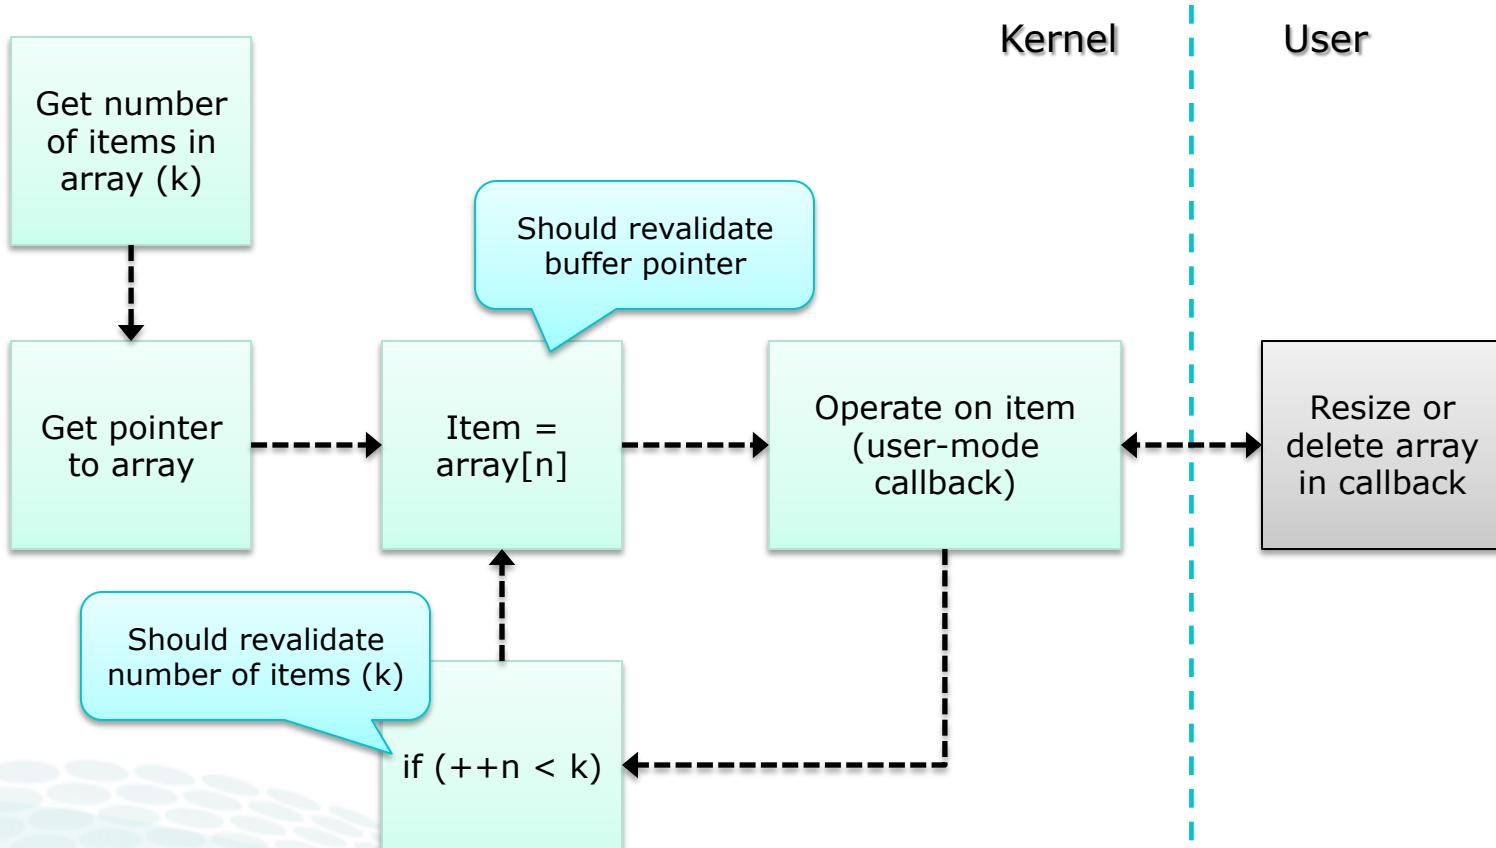
DDE Conversation Object NULL Dereference



Buffer Reallocation

- Many user objects have item arrays or other forms of buffers associated with them
 - E.g. menu items array
- Item arrays where elements are added or removed are often resized to conserve memory
 - Buffer freed if the array is empty
 - Buffer reallocated if elements is above or below a certain threshold
- Any buffer that can be reallocated or freed during a callback must be checked upon return
 - Failure to do so could result in use-after-free

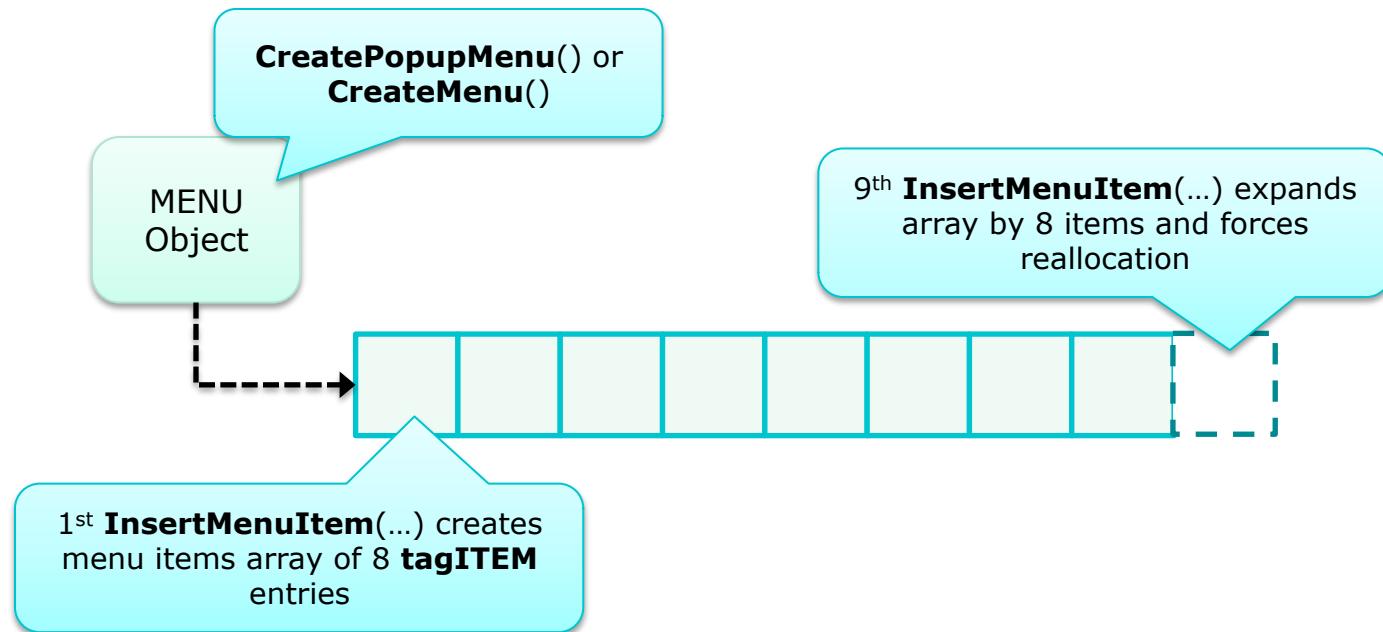
Buffer Reallocation



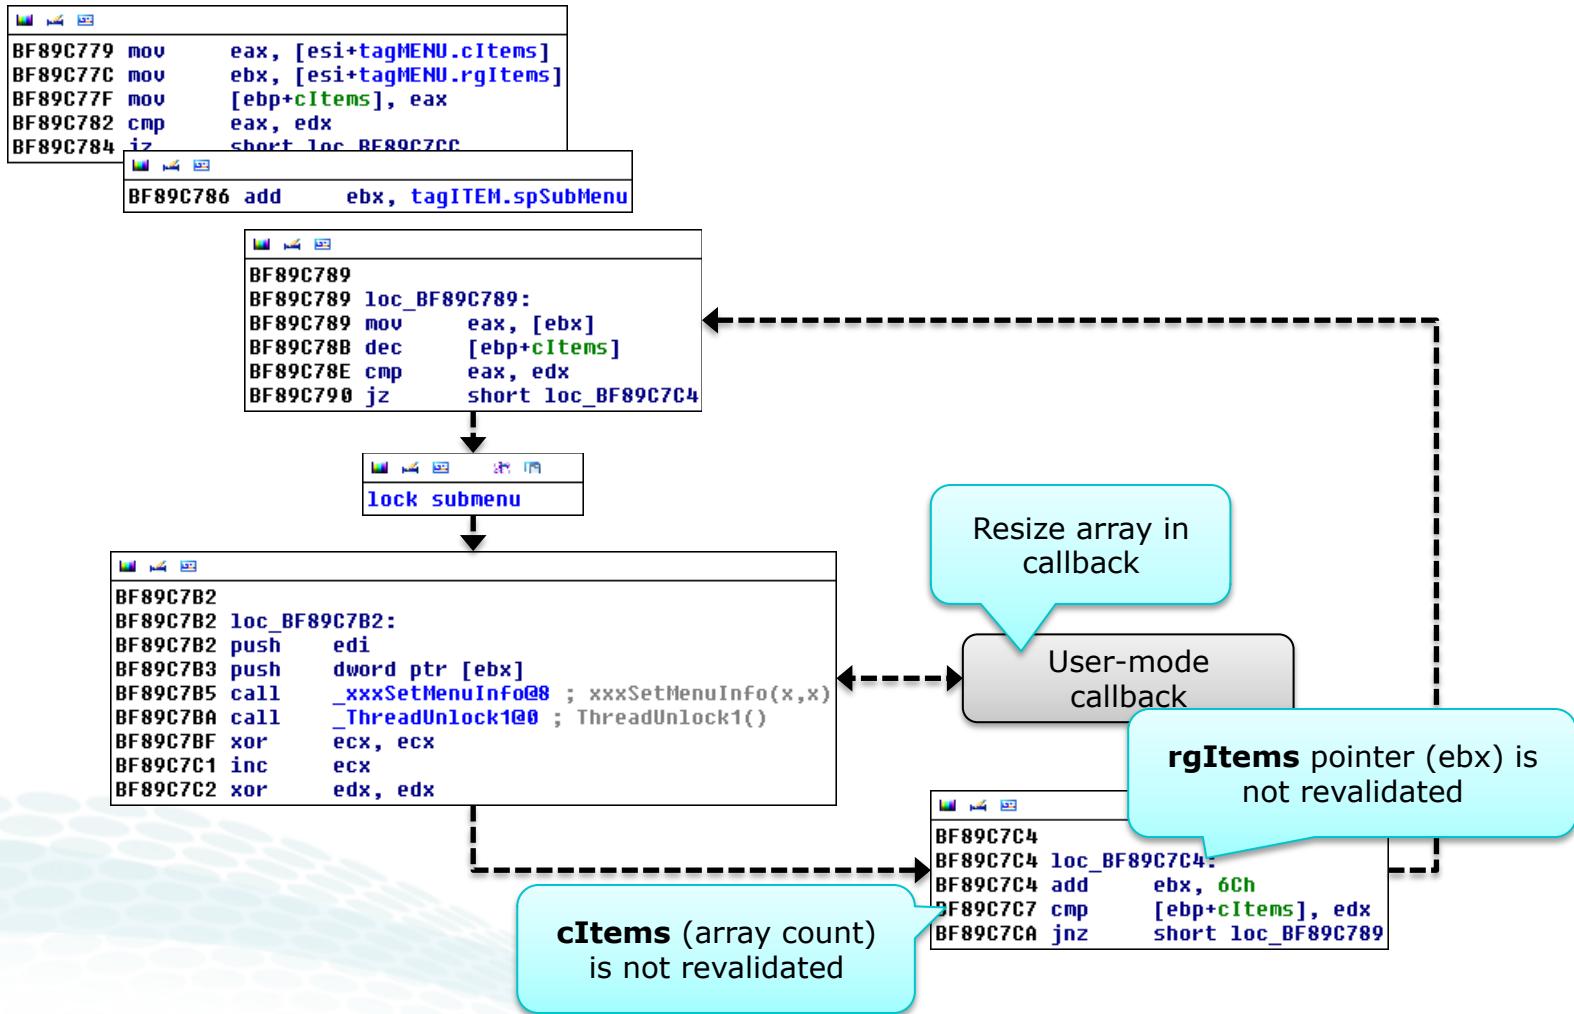
Menu Item Array Use-After-Frees

- Menus may hold an arbitrary number of menu items
 - Stored in a dynamically sized array pointed to by the menu object structure (win32k!tagMENU)
- Win32k did not revalidate the menu items array pointer after user-mode callbacks
 - No way to “lock” a menu item
 - Any ‘xxx’ function operating on menu items was potentially vulnerable
- An attacker could cause the buffer to be reallocated in a callback and trigger a use-after-free

Menu Item Array Reallocation



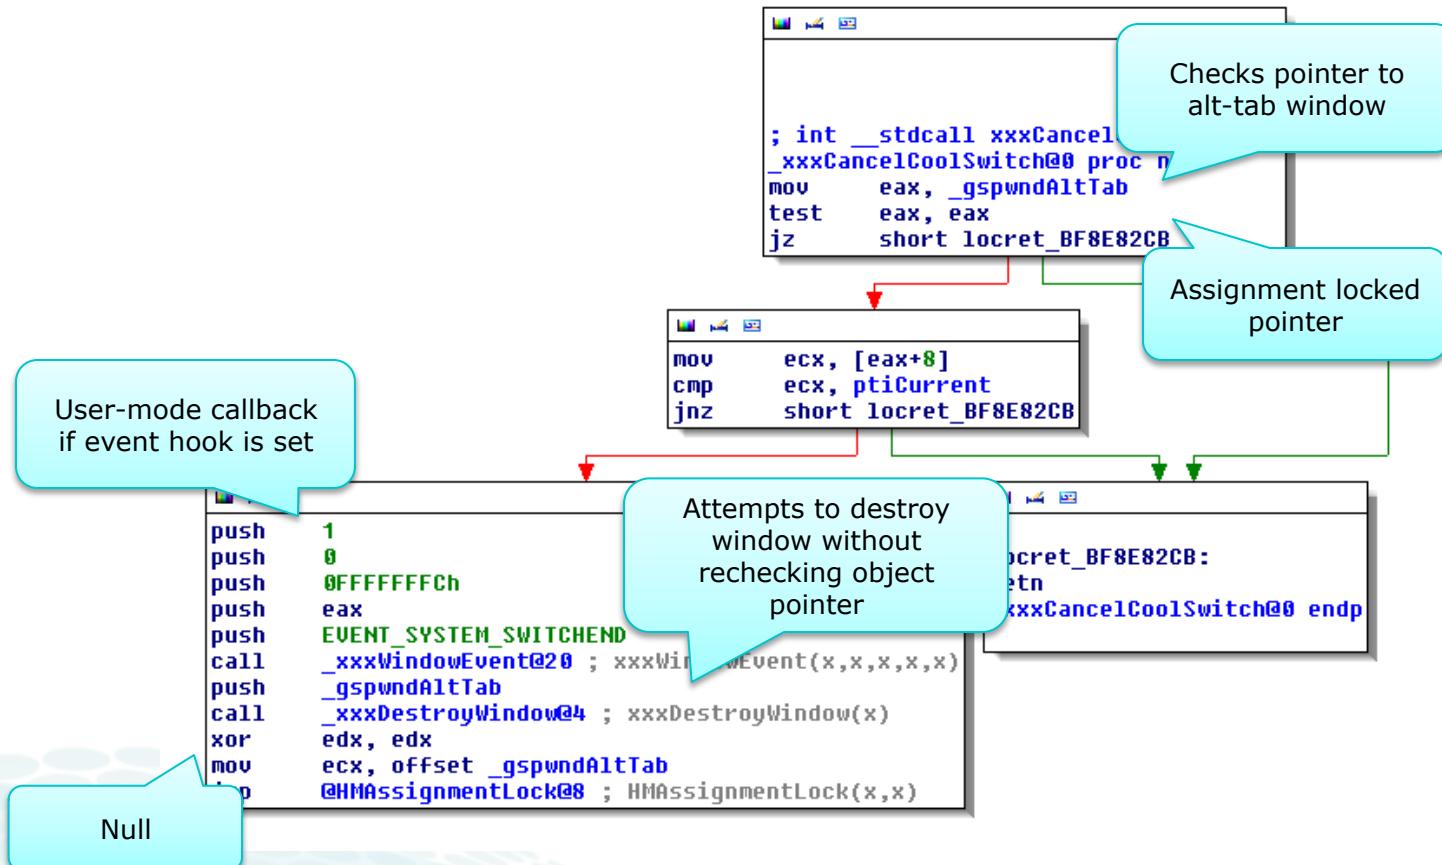
Menu Item Processing Use-After-Free



Time-of-Check-to-Time-of-Use

- The user critical section is generally used to prevent TOCTTOU issues in user object handling
 - User-mode callbacks may allow an attacker to manipulate an object or global value before it is used
- Can be particularly dangerous in clean up routines
 - May invoke callbacks after checks have been made
 - Could result in stale references to objects or buffers
- Values that may have changed must always be (re)checked after a callback has taken place

Time-of-Check-to-Time-of-Use



Handle Validation

- Required to validate handles, their type, and retrieve the corresponding object pointers
 - HMValidateHandle() and friends
- Generic handle validation should be avoided unless the structure of the object is irrelevant
 - Only checks handle table entry and ignores type
- Functions that revalidate handles after callbacks, may no longer be operating on the same object
 - The uniqueness counter designed to provide handle entropy is only 16-bit

Insufficient Handle Validation

```
xxxGetMenuBarInfo(x,x,x,x)+297 mov edi, [ebp+wnd]
xxxGetMenuBarInfo(x,x,x,x)+29A push ecx
xxxGetMenuBarInfo(x,x,x,x)+29B push ecx
xxxGetMenuBarInfo(x,x,x,x)+29C push 1E1h
xxxGetMenuBarInfo(x,x,x,x)+2A1 push edi
xxxGetMenuBarInfo(x,x,x,x)+2A2 call _xxxSendMessage@16 ; xxxSendMessage(x,x,x,x)
xxxGetMenuBarInfo(x,x,x,x)+2A7 mov edx, _ahelist ; handle table pointer
xxxGetMenuBarInfo(x,x,x,x)+2AD mov ecx, eax ; eax: user handle
xxxGetMenuBarInfo(x,x,x,x)+2AF and ecx, 0FFFFh
xxxGetMenuBarInfo(x,x,x,x)+2B5 lea ecx, [ecx+ecx*2]
xxxGetMenuBarInfo(x,x,x,x)+2B8 mov ecx, [edx+ecx*4]
xxxGetMenuBarInfo(x,x,x,x)+2BB test ecx, ecx ; ecx: object pointer
xxxGetMenuBarInfo(x,x,x,x)+2BD jz loc_BF91E14B
```

Function did not check handle type nor validate index in handle table

Function did not check that object was an image (icon/cursor)

```
xxxClientLoadImage(x,x,x,x,x,x,x,x)+16D or dl, 0FFh ; TYPE_GENERIC
xxxClientLoadImage(x,x,x,x,x,x,x,x)+170 mov ecx, edi ; handle from user-mode callback
xxxClientLoadImage(x,x,x,x,x,x,x,x)+172 call @HMValidateHandleNoRip@8 ; HMValidateHandleNoRip(x,x)
xxxClientLoadImage(x,x,x,x,x,x,x,x)+177 mov edi, eax
```

Exploitability

Use-After-Frees and NULL Pointer
Dereferences

Vulnerability Primitives

- Mainly dealing with two vulnerability primitives
 - Use-After-Frees
 - Null-Pointer Dereferences
- Exploitability may depend on the attacker's ability to manipulate heap and pool memory
 - Kernel Pool Exploitation on Windows 7 (BH DC '11)
 - Not much public information on the kernel heap
- Hooking user-mode callbacks is easy
 - NtCurrentPeb()->KernelCallbackTable

Kernel Heap

- The kernel has a stripped down version of the user-mode heap allocator
 - nt!RtlAllocateHeap, nt!RtlFreeHeap, etc.
 - Used by the shared and desktop heaps
- Neither heaps employ any front end allocators
 - ExtendedLookup == NULL
 - No low fragmentation heap or lookaside lists
- Neither heaps encode or obfuscate heap management structures
 - HEAP.EncodeFlagMask == 0

Desktop Heap Base

The screenshot shows the WinDbg debugger interface with the command window displaying the structure of the Desktop Heap Base. The structure is defined as follows:

```
+0x048 CompatibilityFlags : 0
+0x04c EncodeFlagMask : 0
+0x050 Encoding : HEAP_ENTRY
+0x058 PointerKey : 0
+0x05c Interceptor : 0
+0x060 VirtualMemoryThreshold : 0xfe00
+0x064 Signature : 0xffffffff
+0x068 SegmentReserve : 0x100000
+0x06c SegmentCommit : 0x2000
+0x070 DeCommitFreeBlockThreshold : 0x200
+0x074 DeCommitTotalFreeThreshold : 0x2000
+0x078 TotalFreeSize : 0x5a9
+0x07c MaximumAllocationSize : 0x7f
+0x080 ProcessHeapsListIndex : 0
+0x082 HeaderValidateLength : 0x138
+0x084 HeaderValidateCopy : (null)
+0x088 NextAvailableTagIndex : 0
+0x08a MaximumTagIndex : 0
+0x08c TagEntries : (null)
+0x090 UCRList : LIST_ENTRY [ 0xfea14008 - 0xfea1d980 ]
+0x098 AlignRound : 0xf
+0x09c AlignMask : 0xffffffff
+0xa0 VirtualAllocdBlocks : LIST_ENTRY [ 0xteauUUuAU - 0xteauUUuAU ]
+0xa8 SegmentList : LIST_ENTRY [ 0xfea0010 - 0xfea00010 ]
+0xb0 AllocatorBackTraceIndex : 0
+0xb4 NonDedicatedListLength : 0
+0xb8 BlocksIndex : 0xfea00138 Void
+0xc0 UCRIndex : (null)
+0xc4 PseudoTagEntries : (null)
+0xc4 FreeLists : LIST_ENTRY [ 0xfea14008 - 0xfea1d980 ]
+0xc8 LockVariable : (null)
+0xd0 CommitRoutine : 0x93bb69e9    long win32k!UserCommitDesktopMemory+0
+0xd4 FrontEndHeap : (null)
+0xd8 FrontHeapLockCount : 0
+0xda FrontEndHeapType : 0
+0xdc Counters : _HEAP_COUNTERS
```

Annotations with callouts point to specific fields:

- An annotation labeled "EncodingFlagMask and PointerKey" points to the `EncodeFlagMask` and `PointerKey` fields.
- An annotation labeled "No front end allocators" points to the `FrontEndHeap` field.
- An annotation labeled "Free list" points to the `FreeLists` field.
- An annotation labeled "Commit routine to extend the heap" points to the `CommitRoutine` field.

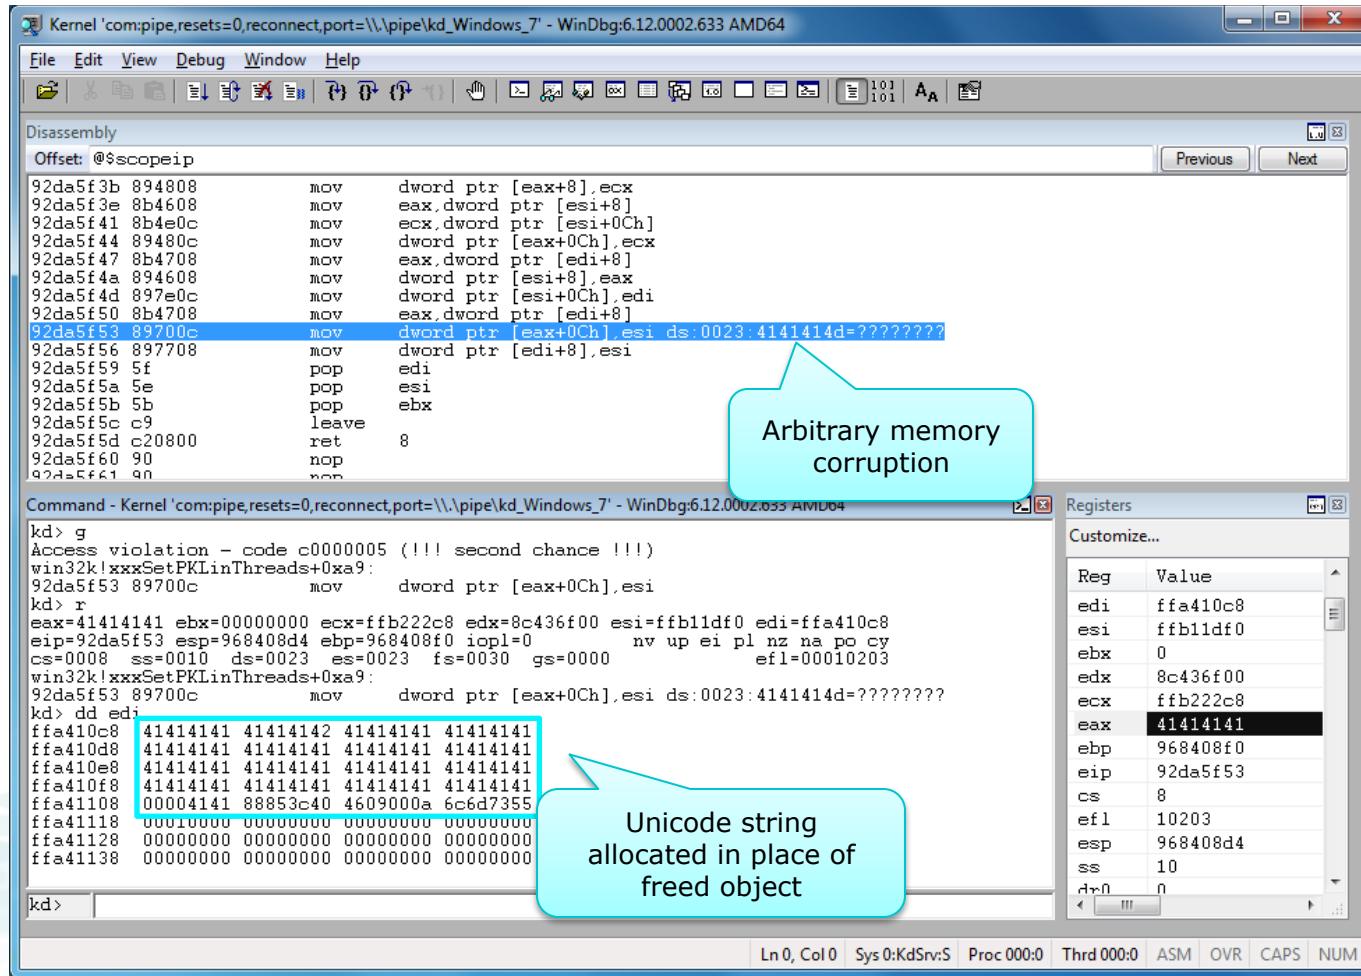
Kernel Heap Management

- Freed memory is indexed into a single free list
 - Ordered by block size
 - *ListHints* used to optimize list lookup
- Requested memory is always pulled from the front of an oversized heap chunk
 - Remaining fragment is put back into the free list
- If the heap runs out of committed memory, win32k calls the *CommitRoutine* to extend the heap
 - Attempts to commit memory from the reserved range
 - E.g. win32k reserves 0xC00000 bytes by default (adjustable by user) for desktop heaps

Use-After-Free Exploitation

- Unicode strings can be used to reallocate freed memory from within user-mode callbacks
 - Allows control of the contents and size of the heap block
 - Caveat: Cannot use WORD NULLs and last two bytes must be NULL to terminate the string
- Desktop heap
 - SetWindowTextW (hWnd, String) ;
- Session pool
 - SetClassLongPtr (hWnd, GCLP_MENUNAME, (LONG) String) ;

Strings As User Objects



Exploiting Object Locking Behavior

- Embedded object pointers in the freed object may allow an attacker to increment (lock) or decrement (unlock) an arbitrary address
 - Common behavior of locking routines
- Some targets
 - HANDLEENTRY.bType
 - Decrement the type of a window handle table entry (1)
 - Destroy routine for free type (0) is null (mappable by user)
 - KAPC.ApcMode
 - Execute code with kernel-mode privileges by decrementing UserMode (1) to KernelMode (0)

Exploiting Object Locking Behavior

Kernel 'com:pipe, resets=0, reconnect, port=\pipe\kd_Windows_7' - WinDbg:6.12.0002.633 AMD64

File Edit View Debug Window Help

Disassembly

Offset: @\$scopeip

```
8216c555 90          nop
win32k!HMUnlockObject:
8216c556 8bff         mov    edi,edi
8216c558 55          push   ebp
8216c559 8bec         mov    ebp,esp
8216c55b 8b4508       mov    eax,dword ptr [ebp+8]
8216c55e ff4804       dec    dword ptr [eax+4]    ds:0023:deadbeef=?????????
8216c561 7506         jne    win32k!HMUnlockObject+0x13 (8216c569)
8216c563 50          push   eax
8216c564 e8cf3a0000  call   win32k!HMUnlockObjectInternal (82170038)
8216c569 5d          pop    ebp
8216c56a c20400       ret    4
8216c56d 90          nop
```

Unlocking user-controlled pointer (0xdeadbeef)

Command - Kernel 'com:pipe, resets=0, reconnect, port=\pipe\kd_Windows_7' - WinDbg:6.12.0002.633 AMD64

```
kd> r
eax=deadbeeb ebx=fe95a990 ecx=ff910000 edx=fea11480 esi=fea11480 edi=deadbeeb
eip=8216c55e esp=9431dc00 ebp=9431dc00 iopl=0 nv up ei ng nz na pe nc
cs=0008 ss=0010 ds=0023 es=0023 fs=0030 gs=0000 efl=000000286
win32k!HMUnlockObject+0x8:
8216c55e ff4804       dec    dword ptr [eax+4]    ds:0023:deadbeef=?????????
```

kd> kb
ChildEBP RetAddr Args to Child
9431dc00 8216c9e0 deadbeeb 00000000 fe95a978 win32k!HMUnlockObject+0x8
9431dc00 820d0cb1 820d0b8b 0048fa0c 002af85c win32k!HMAssignmentLock+0x45
9431dc00 820d0bb3 9431dcfc 9431dcf8 9431dcf4 win32k!xxxCsDdeInitialize+0x67
9431dd18 8285542a 0048fa0c 0048fa1c 0048fa14 win32k!NtUserDdeInitialize+0x28
9431dd18 779464f4 0048fa0c 0048fa1c 0048fa14 nt!KiFastCallEntry+0x12a
002af91c 7795b3f5 7ffd000 77a4624b 00000000 ntdll!KiFastSystemCallRet
002af95c 7795b3c8 00fe16e2 7ffd000 00000000 ntdll!_RtlUserThreadStart+0x70
002af974 00000000 00fe16e2 7ffd000 00000000 ntdll!_RtlUserThreadStart+0x1b

HMAssignmentLock unlocks the existing user-controlled pointer

Ln 0, Col 0 Sys 0:KdSrv:S Proc 000:0 Thrd 000:0 ASM OVR CAPS NUM

NULL Pointer Vulnerabilities

- Potentially exploitable on the Windows platform
 - Non-privileged users can map the null page, e.g. via NtAllocateVirtualMemory or NtMapViewOfFile
- Many NULL pointer vulnerabilities are concerned with window object pointers
- An attacker could map the null page and set up a fake window object
 - E.g. define a server-side window procedure and handle messages with kernel level privileges

NULL Pointer Object Exploitation

Kernel 'com:pipe, resets=0, reconnect, port=\\.\pipe\kd_WinXP_SP3_dev' - WinDbg:6.12.0002.633 AMD64

File Edit View Debug Window Help

Disassembly

Offset: @\$scopeip

No prior disassembly possible

41414141 ?? ???
41414142 ?? ???
41414143 ?? ???
41414144 ?? ???
41414145 ?? ???

Command - Kernel 'com:pipe, resets=0, reconnect, port=\\.\pipe\kd_WinXP_SP3_dev' - WinDbg:6.12.0002.633 AMD64

kd> r
eax=00004688 ebx=00000002 ecx=f06e0688 edx=00000000 esi=00000000 edi=e105d830
eip=41414141 esp=f06e0634 ebp=f06e0670 iopl=0 nv up ei pl zr na
cs=0008 ss=0010 ds=0023 es=0023 fs=0030 gs=0000
41414141 ?? ???
kd> kb 6
ChildEBP RetAddr Args to Child
WARNING: Frame IP not in any known module. Following frames may be wrong.
f06e0630 bf814099 00000000 00000002 00000000 0x41414141
f06e0670 bf80ecc6 00000000 00000002 00000000 win32k!xxxSendMessageTimeout+0x18a
f06e0694 bf8457c1 00000000 00000002 00000000 win32k!xxxSendMessage+0x1b
f06e06d0 bf845eff 00000000 00000000 e105d830 win32k!xxxDW_SendDestroyMessages+0x35
f06e071c bf8e82bf 00000000 bf91e8fa bc69eee0 win32k!xxxDestroyWindow+0x381
f06e0724 bf91e8fa bc69eee0 00000000 0000f040 win32k!xxxCancelCoolSwitch+0x2d

kd> dd 0
00000000 000200ac 00000000 e105d830 00000000
00000010 00000000 00040000 00000000 00000000
00000020 00000000 00000000 00000000 00000000
00000030 00000000 00000000 00000000 00000000
00000040 00000000 00000000 00000000 00000000
00000050 00000000 00000000 00000000 00000000
00000060 41414141 00000000 00000000 00000000
00000070 00000000 00000000 00000000 00000000

kd>

Ln 0, Col 0 Sys 0:KdSrv:S Proc 000:0 Thrd 000:0 ASM OVR CAPS NUM

Message sent to null pointer object

Fake null page window object

Server-side window procedure pointer

Demo

- **Window Object Use-After-Free
(CVE-2011-1237)**
 - Arbitrary kernel code execution via HANDLEENTRY corruption

Mitigations

Protecting Against Privilege Escalation
Vulnerabilities

Mitigating Use-After-Free Exploitation

- Need to address an attacker's ability to reallocate the freed memory before use
- Some approaches
 - Delayed frees while processing a callback
 - Dedicated free lists for user objects
 - Isolate strings used in reallocating memory
 - Track allocations between ring transitions, e.g. pointers on the stack before a callback
- Generally hard to mitigate without significantly impacting performance

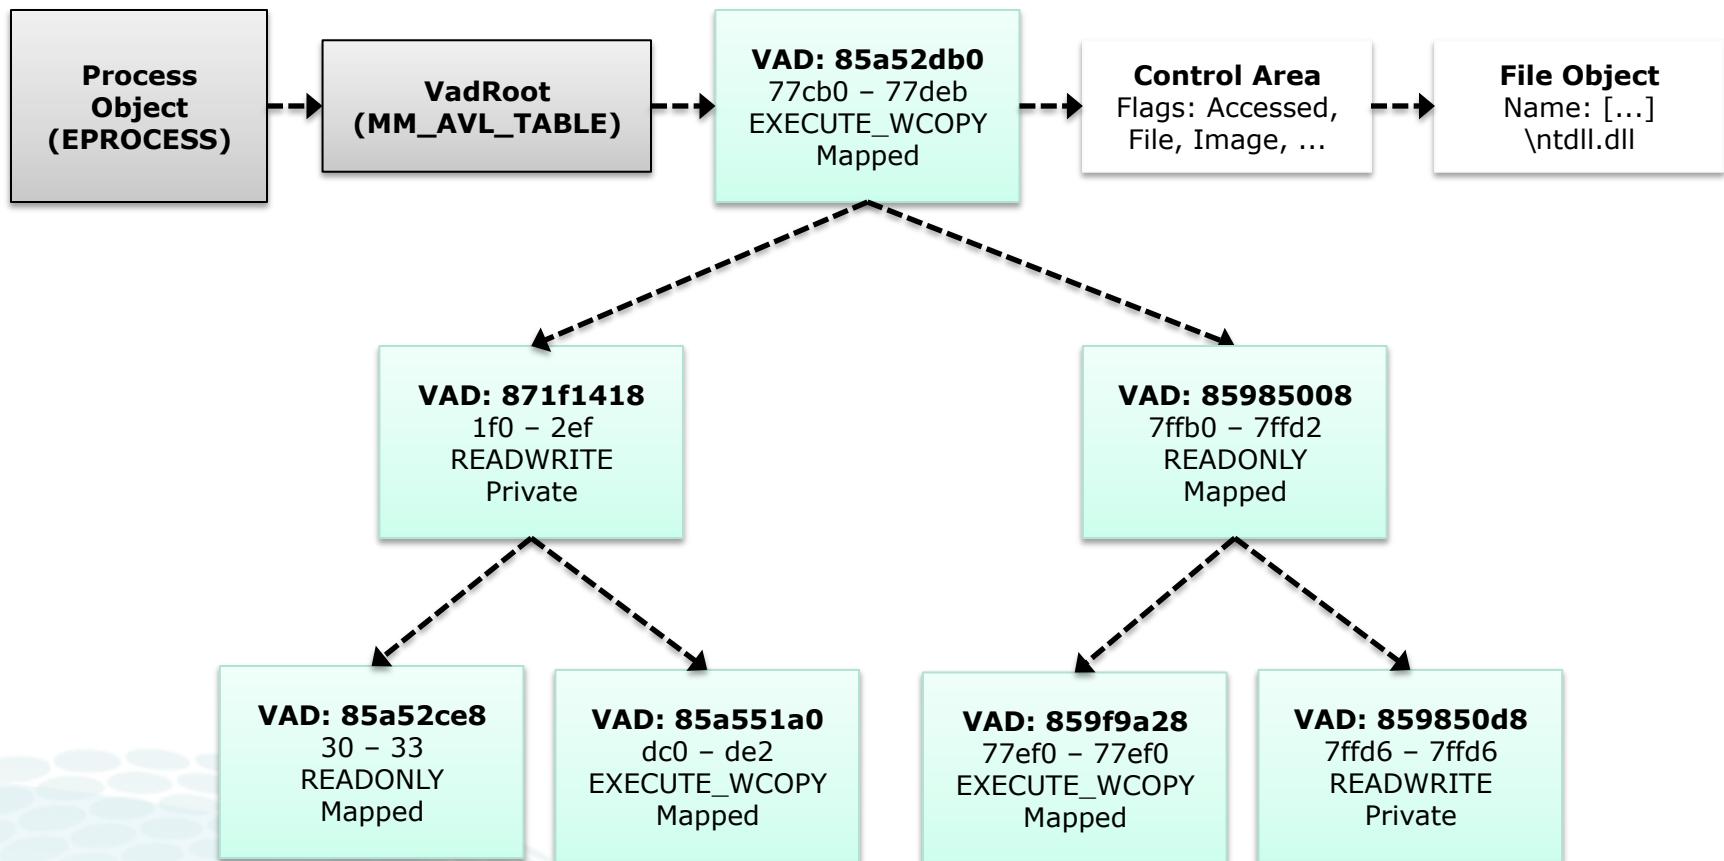
Mitigating NULL Pointer Exploitation

- We can address null pointer exploitation by denying users the ability to map the null page
- Some potential ways of addressing null page mappings
 - System call hooking
 - Page Table Entry (PTE) modification
 - VAD manipulation
- System call hooking not supported on x64
- PTE modification requires page to be mapped

VAD Manipulation

- User mode process space is described using **Virtual Address Descriptors (VADs)**
 - Structured in self-balanced AVL trees
- **VADs are always checked before PTEs are created**
 - E.g. used to implement the NO_ACCESS protection
- **VADs are used to secure memory, e.g. can be made non-deletable**
 - PEBS and TEBs
 - KUSER_SHARED_DATA section

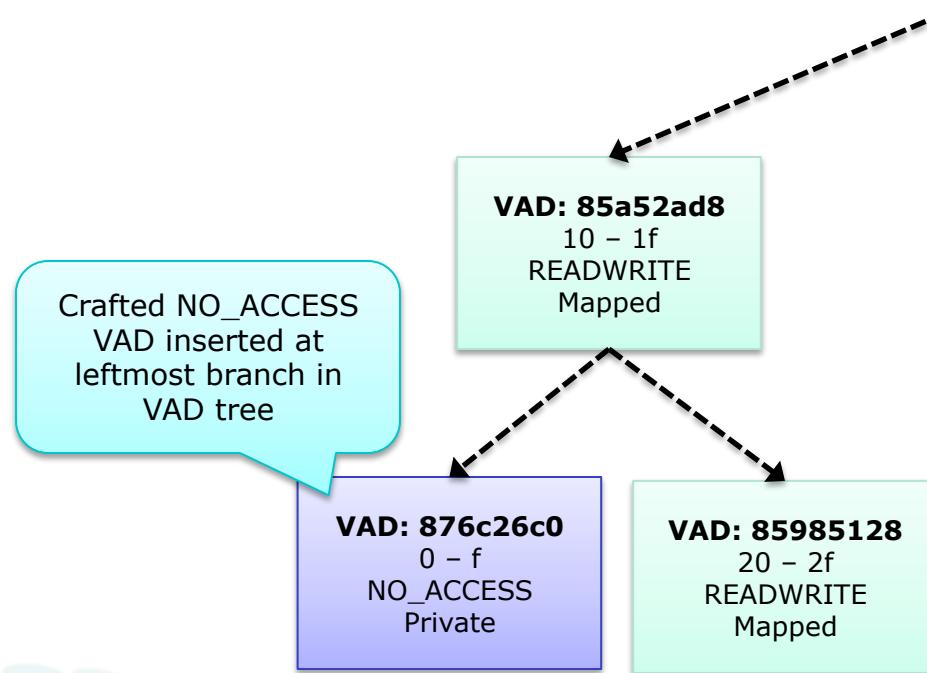
VAD Tree



Restricting Null Page Access

- We insert a crafted VAD entry to restrict null page access
 - Ring3 code cannot modify the VAD entry
- Avoid deletion using the same method employed by PEBs and TEBs
 - Secure address range from 0 up to 0xFFFF
 - Set protection to NO_ACCESS
- Use a special VAD flag to prevent memory commits
 - Protection cannot be changed on uncommitted memory!

VAD Tree /w Crafted Entry



Manipulated Process VAD Tree

The screenshot shows a WinDbg session with the following command:

```
kd> !process @$proc 0
```

Output:

```
PROCESS 85a61ab8 SessionId: 1 Cid: 0eec Peb: 7ffd9000 Pa  
DirBase: 1f05b3e0 ObjectTable: 962db880 HandleCount: 6  
Image: ktest.exe
```

Another command:

```
kd> r? $t0 = &((nt!_EPROCESS *)0)->VadRoot;!vad poi(@$proc+$t0+8);
```

Output:

VAD	level	start	end	commit	Type	Description
859ded48	(4)	0	f	524287	Private	NO_ACCESS
85a7f0d8	(3)	10	1f	0	Mapped	READWRITE
85aeee360	(4)	20	2f	0	Mapped	READWRITE
866007f8	(2)	30	33	0	Mapped	READONLY
877f5e30	(3)	40	40	1	Private	READWRITE
859aa0b8	(4)	50	b6	0	Mapped	READONLY
875c0390	(1)	1b0	2af	3	Private	READWRITE
8778c140	(4)	400	40f	3	Private	READWRITE
8719c0b0	(3)	430	52f	15	Private	READWRITE
85a52eb8	(2)	1020	1042	6	Mapped	Exe EXECUTE_WRITECOPY
859f92f0	(4)	75f60	75fa9	3	Mapped	Exe EXECUTE_WRITECOPY
859b2978	(3)	76e80	76f53	2	Mapped	Exe EXECUTE_WRITECOPY
877d66f0	(0)	77cb0	77deb	9	Mapped	Exe EXECUTE_WRITECOPY
8756d630	(2)	77ef0	77ef0	0	Mapped	Exe EXECUTE_WRITECOPY
877de058	(3)	7f6f0	7f7ef	0	Mapped	READONLY
85a97078	(1)	7ffb0	7ffd2	0	Mapped	READONLY
86681b88	(2)	7ffd9	7ffd9	1	Private	READWRITE
86185e28	(3)	7ffdf	7ffdf	1	Private	READWRITE

Total VADs: 18 average level: 3 maximum depth: 4

Another command:

```
kd> !pte 0
```

Output:

```
VA 00000000  
PDE at C0600000 PTE at C0000000  
contains 0000000006E3E867 contains 0000000000000000  
pfn 6e3e ---DA--UWEV not valid
```

A callout bubble points to the first VAD entry:

Crafted NO_ACCESS VAD

A callout bubble points to the bottom of the memory dump:

Invalid memory

Bottom status bar:

```
Ln 0, Col 0 Sys 0:KdSrv:S Proc 000:0 Thrd 000:0 ASM OVR CAPS NUM
```

Mitigation Results

Function	Addr	Type	Protection	Result
NtAllocateVirtualMemory	1	MEM_RESERVE	READONLY	0xC0000018
NtAllocateVirtualMemory	1	MEM_COMMIT	READONLY	0xC0000018
NtMapViewOfSection	1	MEM_DOS_LIM*	READONLY	0xC0000018
NtProtectVirtualMemory	0		READWRITE	0xC000002D
NtProtectVirtualmemory	0		READONLY	0xC0000045
NtFreeVirtualMemory	0	MEM_RELEASE		0xC0000045

0xC0000018	STATUS_CONFLICTING_ADDRESSES
0xC000002D	STATUS_NOT_COMMITTED
0xC0000045	STATUS_INVALID_PAGE_PROTECTION

*Allows section mapping on page boundary on x86 platforms

Demo

- Null page mapping mitigation

Conclusion

Remarks and Conclusion

Future of the Win32k Subsystem

- **Win32k needs a much more consistent and security oriented design**
 - It should not be necessary for the kernel to make direct calls back into user-mode
 - Reconsider performance benefit of shared user and kernel-mode memory mappings
- **The Window Manager should provide mutual exclusion on a per-object basis**
 - Better suited towards multicore architectures
 - Similar to what is done in GDI and the NT executive

Conclusion

- Legacy components constitute the most vulnerable parts of an operating system
 - Security is not usually part of the original design
 - Win32k is built around very old GUI subsystem code
- Kernel exploitation requires knowledge about the kernel address space
 - Limiting access to such information is important
- Although hard, mitigating Windows kernel exploitation is possible

References

- **Windows Kernel Internals: Win32K.sys**
 - David B. Probert, Microsoft
- **Analyzing Local Privilege Escalations in win32k**
 - mxatone (Uninformed Vol. 10)
- **Windows Creation Vulnerability (MS10-048)**
 - Nicolás Economou
- **Pointers and Handles: A Story of Unchecked Assumptions in the Windows Kernel**
 - Alex Ionescu
- **Understanding the Low Fragmentation Heap**
 - Chris Valasek

Questions ?

- Email: kernelpool@gmail.com
- Blog: <http://mista.nu/blog>
- Twitter: @kernelpool
- Norman MDT Blog:
<http://blogs.norman.com/category/malware-detection-team>